

Linux Kernel Memory Ordering Model: Concepts, Theories and Tools

Boqun Feng

Linux Technology Center, IBM

Self-Introduction

- 冯博群(Boqun Feng)
- Join IBM LTC at 2014 April, Working on Linux Kernel since then
- Focus on scheduler, locking, RCU and other core kernel areas

When We Talk about Memory Ordering

- What We Talk About?

*A: Hey, We need you to
work on memory ordering*

*B: OK! No Problem. How
much memory we need to
order?*

A Real World Ordering Problem

- “内核和应用同时访问一个物理内存，比如大小为8字节,内核需要更新这个8字节内存对应的内容，应用去读这8字节物理内存里面的内容;反过来应用有时也会去写，内核去读。有时写完后，去读时有时读取为-1，有时可能是其他值。尝试多次重新读取后，正确了。”
- Parallel Program and Shared Memory

Agenda

- Parallel Programming in Linux Kernel: Two Facts and One Question
- Primitives and Semantics
- The role of cycles
- Litmus Tests and Tools

Fact #1

- **More and More Complex but “Efficient” Parallel Programming Algorithms are getting into Linux Kernel.**
 - Filesystem: Pathname lookup rewrite: <https://lwn.net/Articles/649115/>
 - Network: Lockless TCP Listener: <https://lwn.net/Articles/659199/>
 - Even Locking Primitives: Qspinlock: <https://lwn.net/Articles/590243/>
- Why “Efficient”
 - Parallel programs are not definitely efficient
 - Parallel programs are not easily to be correct
- But Incorrect parallel programs are usually efficient.

Fact #2

- **Yet We Still Could Find Parallel Programming Bugs Existed in Linux Kernel For a Long Time**
 - spin_unlock_wait() on PPC and ARM
 - <https://marc.info/?l=linux-kernel&m=144731258921696>
 - spin_unlock_wait() for qspinlock on x86
 - <https://marc.info/?l=linux-kernel&m=146372279722288>
 - Race in try_to_wake_up()
 - <https://marc.info/?l=linux-kernel&m=147263879404805>

One Question

- How Could We Verify Our Parallel Code?
 - To ensure new “efficient” algorithms are correct
 - To find old bugs in kernel

Answer

- Before Year 2014, we had:
 - Documentation/memory-barriers.txt, the *Children-Frightener*.
 - Several people who could review patches for memory ordering issues.
 - Paul Mckenney, Peter Zijlstra, Will Deacon, etc.
- Around Year 2014 and 2015, things are better, because we had more:
 - Me ;-)
 - Several researches and tools for arch-specific memory ordering models:
 - PPCMEM and ARMMEM, and related researches
 - Herd and related researches
- Now and Future, we are going to have:
 - A Formal Ordering Model:
 - LCE 2016: *Linux-Kernel Memory Ordering: Help Arrives At Last!*
 - by Paul E. Mckenney, Jade Alglave, Luc Maranget, Andrea Parri, and Alan Stern and more

When We Talk about Memory Ordering

- What We Talk About?
 - Things that could reorder memory operations
 - Things that could prevent reordering, IOW, could order memory operations.

Things That Could Reorder Memory Ops

- Compiler
 - "A compiler is within its right to" ...
 - *N4455: No Sane Compiler Would Optimize Atomics*
- Hardware
 - Out-of-order execution
 - Store buffer and invalidate queue.

From a viewpoint of a memory ordering model, nothing is ordered, unless the model says yes.

Things That Order Memory Operations

- Barriers
 - Compiler Barrier: `barrier()`
 - Full barrier: `mfence(x86)`, `sync(PPC)`
 - Partial Barrier: `lwsync`
 - Operation with Barrier Semantics: `ldaxr`, `stlxr(AARCH64)`
- Dependencies
- Arch-dependent Intrinsic Ordering
 - TSO, SC

Based on which, we build Linux internal primitives on ordering.

Agenda

- Parallel Programming in Linux Kernel: Two Facts and One Question
- **Primitives and Semantics**
- The role of cycles
- Litmus Tests and Tools

Relationship Between Memory Ops

- Ordering is one kind of relationship on the set of memory operations.
- And ordering primitives and concepts will make some relationships special
 - i.e. Make a relationship provide ordering (in some cases).
- Let's go through some kernel primitives relationship-wisely.


Linux Primitives for Memory Ordering

- `smp_mb()`
- `smp_wmb()` and `smp_rmb()`
- `smp_load_acquire()` and `smp_store_release()`
- `lockless_dereference()`, `rcu_dereference()` and `rcu_assign_pointer()`

smp_mb()

- A full barrier
 - Orders READ->READ, READ->WRITE, WRITE->READ, WRITE->WRITE
 - Provides Transitivity
 - Implemented as mfence on x86_64, sync on PPC

```
WRITE_ONCE(*x, 1);  
  
smp_mb();  
  
r1 = READ_ONCE(*y, 1);
```




smp_wmb() and smp_rmb()

- Partial barriers:
 - smp_wmb() orders WRITE->WRITE
 - smp_rmb() orders READ->READ
 - Implemented as lwsync on PPC

```
WRITE_ONCE(*x, 1);
```

```
smp_wmb();
```

```
WRITE_ONCE(*y, 1);
```



smp_load_acquire() and smp_store_release()

- smp_load_acquire()
 - A load
 - Every operation (program-order) after smp_load_acquire() will not happen before load.
 - One-way barrier
- smp_store_release()
 - A store
 - Every operation (program-order) before smp_store_release() will not happen after the store
 - One-way barrier

smp_load_acquire() and smp_store_release()

```
WRITE_ONCE(*a, 1)


r1 = smp_load_acquire(x);

WRITE_ONCE(*y, 1);

r2 = READ_ONCE(*z);

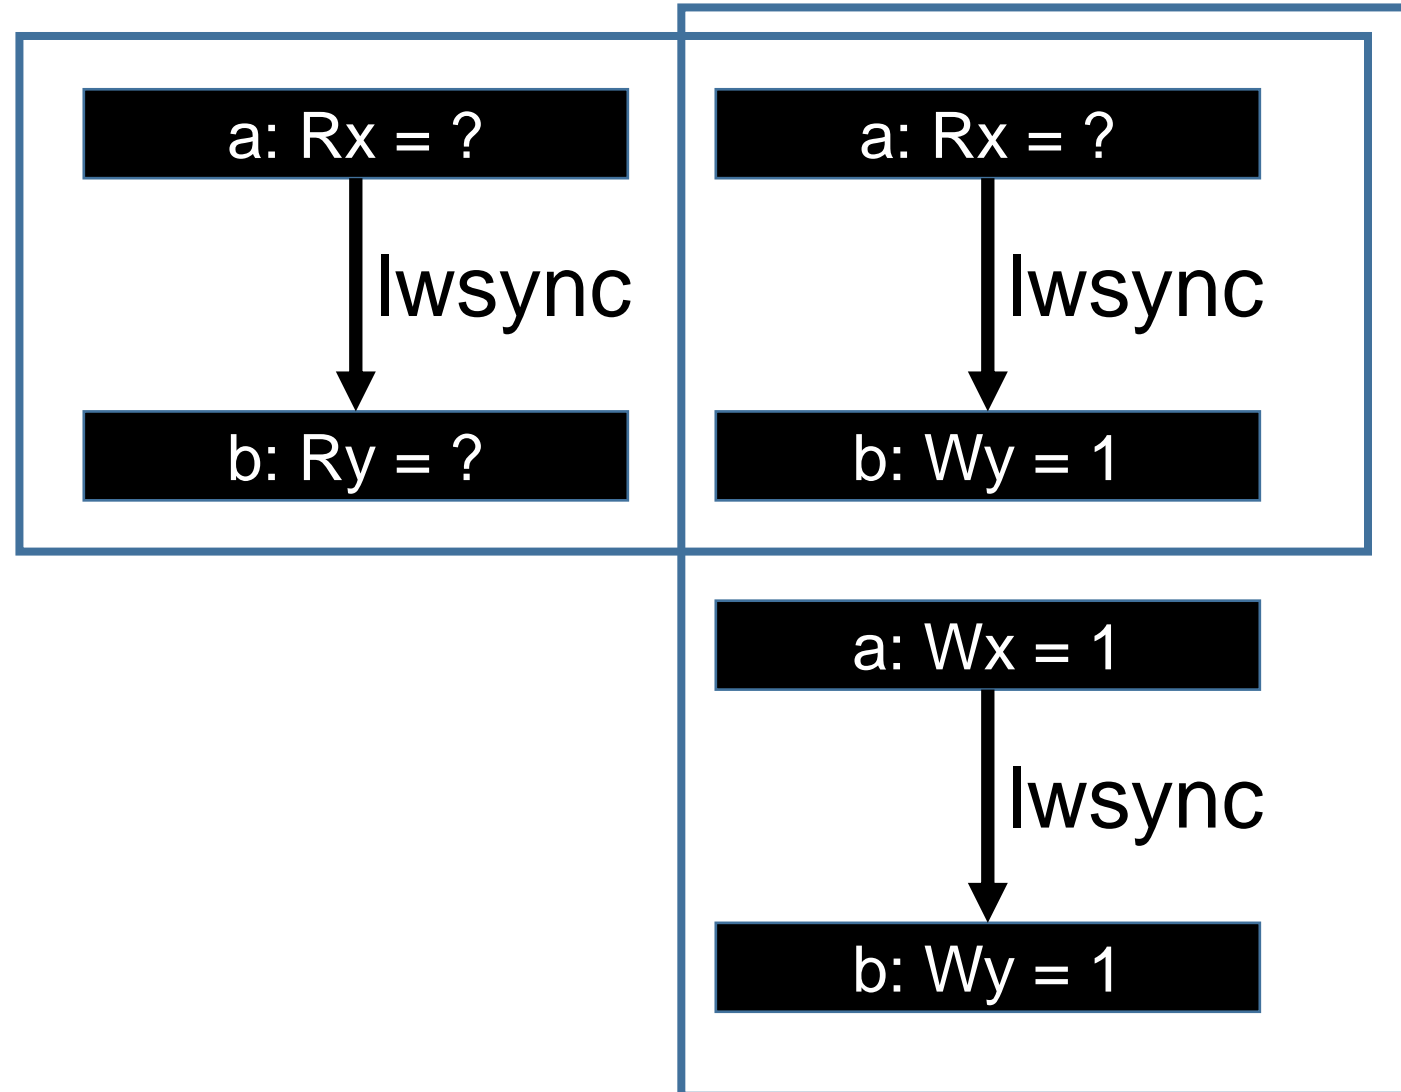
smp_store_release(x, 1);

r3 = READ_ONCE(*a);
```



ACQUIRE and RELEASE: Implementation

- lwsync on PPC
 - Orders READ->READ
 - Orders READ->WRITE
 - Orders WRITE->WRITE
- ACQUIRE: load+lwsync
- RELEASE: lwsync+store



*_dereference()

- lockless_dereference() and rcu_dereference()
 - Provides ordering between memory operations having data/address dependencies.
- rcu_assign_pointer()
 - RELEASE?

Dependencies


- Data/Address
 - READ->{READ, WRITE}
- Control
 - READ->WRITE

```
<Data/Address Dependencies>  
struct T *ptr = READ_ONCE(*p);
```

```
tmp = ptr->a; // or ptr->a = 1;
```



```
<Control Dependencies>  
struct T *ptr = READ_ONCE(*p);  
  
if (ptr) {  
    WRITE_ONCE(x, 1);  
    ...  
}
```



a: Rp = ?

addr

b: Ra = ?

a: Rp = ?

ctrl

b: Wx = 1

Orders Provided by Linux Kernel Primitives

- Control Dependencies if no compiler reordering involved
- Data Dependencies if leaded by a `*_dereference()` primitives
- ACQUIRE and RELEASE ordering by their definitions, i.e. one way barrier
- Memory Operations before and after `smp_mb()`
- Reads before and after `smp_rmb()`
- Writes before and after `smp_wmb()`

Agenda

- Parallel Programming in Linux Kernel: Two Facts and One Question
- Primitives and Semantics
- **The role of cycles**
- Litmus Tests and Tools

Communication


- Orders provided on one CPU are useless unless in the context of communication.
- We usually use the word "Pairing" or "Synchronize" to describe an communication.

Ways of Communication

- READ FROM
- COHERENCE
- FROM READ


READ FROM(rf)

- A READ observes the value from a WRITE




COHERENCE (co)

- A WRITE observes another WRITE via overwriting.
- A Total Order Relationship on WRITES to a same variable.



FROM READ(fr)

- A WRITE observes a READ via failing to be observed by the READ



Cycles: Questions

- x and y are zeros initially
- P0 executes on CPU 0
- P1 executes on CPU 1
- if `r1 == 1` will `r2 == 0`?

```
P0(int *x, int *y)
{
    WRITE_ONCE(*x, 1);
    smp_store_release(y, 1);
}

P1(int *x, int *y)
{
    int r1, r2;

    r1 = smp_load_acquire(y);
    r2 = READ_ONCE(*x);
}
```


Cycles

- if $r1 == 1$ will $r2 == 0$?

```
P0(int *x, int *y)
{
    WRITE_ONCE(*x, 1);
    smp_store_release(y, 1);
}


P1(int *x, int *y)
{
    int r1, r2;

    r1 = smp_load_acquire(y);
    r2 = READ_ONCE(*x);
}
```



Cycles

- Cycles of relationships may mean cycles of timings
- Do we allow the following cycle?
- What does it mean if we disallow this cycle?



Cycles for modeling

- Memory ordering model could be defined as what kind of cycles of relationship we should prohibit.
- And cycles are easily to detect by tools

Cycles Prohibits in Linux Kernel(In Informal Words)

- Cycles built by communication relationship(**com**, which consists of **rf**, **fr** and **co**).
- Cycles built by **po-loc**(program order for a memory location, relationship between memory operations on the some variable with program order) or **po**
- Cycles built by either **com** or **po-loc**, this is actually cache coherence.
- Cycles built by at most one **fr** or **rf** or **acq-po** or **po-rel** or other relationship provided by kernel primitives, whose corresponding program pattern fulfills maybe 80% daily use.
- For more? See the formal model.

Agenda

- Parallel Programming in Linux Kernel: Two Facts and One Question
- Primitives and Semantics
- The role of cycles
- **Litmus Tests and Tools**

Litmus Tests

- Pseudo parallel code fragments
- A baby step to verify a whole project
 - Useful for discussions
 - Easy to name and check

```
{ x = 0; y = 0}
P0(int *x, int *y)
{
    WRITE_ONCE(*x, 1);
    smp_store_release(y, 1);
}
P1(int *x, int *y)
{
    int r1, r2;

    r1 = smp_load_acquire(y);
    r2 = READ_ONCE(*x);
}
exists
(1:r1 = 1  $\wedge$  1:r2 = 0)
```

Three Parts of Litmus Tests

- Initial Value

- Code

- Exist-clause: An Assertion

```
{ x = 0; y = 0}
```

```
P0(int *x, int *y)
```

```
{
```

```
    WRITE_ONCE(*x, 1);
```

```
    smp_store_release(y, 1);
```

```
}
```

```
P1(int *x, int *y)
```

```
{
```

```
    int r1, r2;
```

```
    r1 = smp_load_acquire(y);
```

```
    r2 = READ_ONCE(*x);
```

```
}
```

```
exists
```

```
(1:r1 = 1  $\wedge$  1:r2 = 0)
```

Litmus Tests Could be Code or Assembly

```
{ x = 0; y = 0}
P0(int *x, int *y)
{
    WRITE_ONCE(*x, 1);
    smp_store_release(y, 1);
}
P1(int *x, int *y)
{
    int r1, r2;

    r1 = smp_load_acquire(y);
    r2 = READ_ONCE(*x);
}
exists
(1:r1 = 1  $\wedge$  1:r2 = 0)
```



```
{
0:r10=x;0:r11=y;0:r3=1
1:r10=x;1:r11=y;
x=0;y=0;
}

P0          | P1          ;
stw r3, 0(r10) | lwz r1, 0(r11) ;
lwsync      | lwsync      ;
stw r3, 0(r11) | lwz r2, 0(r10) ;

exists
(1:r1 = 1  $\wedge$  1:r2 = 0)
```

Litmus Tests Could be Checked by Tools

States 3

1:r1=0; 1:r2=0;

1:r1=0; 1:r2=1;

1:r1=1; 1:r2=1;

No


Witnesses

Positive: 0 Negative: 3

Condition exists $(1:r1=1 \wedge 1:r2=0)$

Observation rel-acq-mp Never 0 3

Never



Tools -- PPCMEM

- PPCMEM
 - An Emulator that could emulate modeled memory events of PPC architectures.
 - Input: a small code snippet called litmus
 - Output: result of all possible executions of the given litmus test, and whether a designed assertion fails or not.
 - Pros: Having an interactive interface for understanding the hardware behavior
 - Cons: Slow and not “portable” for other models.

Tools -- Herd

- Herd
 - Generate all possible execution candidates, and examine each candidate with a given model, leave the ones survive after the examination as the possible result of the model.
 - Input: a litmus test and a memory ordering model(having default models for x86, AARCH, PPC, etc.)
 - Output: result of all possible executions of the given litmus test, and whether a designed assertion fails or not.
 - Pros: Fast, and allow self-defined models
 - Cons: Don't have a interactive interface.

Conclusion

- We are going to have more and more “fun” in parallel program in kernel
- Luckily after 25 years of development, we have more docs and tools to help us.
- In the future, we will have a more formal model and hopefully, it could solve all the problems in kernel parallel programming ;-)

Q & A

A singleton

Work on a single thread

```
struct T {  
    int a;  
};  
...  
static struct T *instance;  
  
struct T *get_instance()  
{  
    if (instance == NULL) {  
        instance = malloc(sizeof(struct T));  
        instance->a = some_value;  
    }  
    return instance;  
}
```

A multi-thread-safe singleton: first try

```
static struct T *instance;

struct T *get_instance()
{
    if (instance == NULL) {
        instance = malloc(sizeof(struct T));
        instance->a = some_value;
    }
    return instance;
}
```

won't work:

race when get_instance() is called by two thread in the same time.

malloc twice.

A multi-thread-safe singleton

```
static struct T *instance;
static spinlock_t instance_lock;

struct T *get_instance()
{
    if (instance == NULL) {
        spin_lock(&instance_lock);
        if (instance == NULL) {
            instance = malloc(sizeof(struct T));
            instance->a = some_value;
        }
        spin_unlock(&instance_lock);
    }
    return instance;
}
// won't work, may observe an uninitialized value of ->a
```

A multi-thread-safe singleton: simple solution

```
static struct T *instance;  
static spinlock_t instance_lock;  
  
struct T *get_instance()  
{  
    spin_lock(&instance_lock);  
    if (instance == NULL) {  
        instance = malloc(sizeof(struct T));  
        instance->a = some_value;  
    }  
    spin_unlock(&instance_lock);  
    return instance;  
}
```

Simple solution:
Adding a lock.

**Work, but need to acquire
the lock every time**

A multi-thread-safe singleton: try to optimize

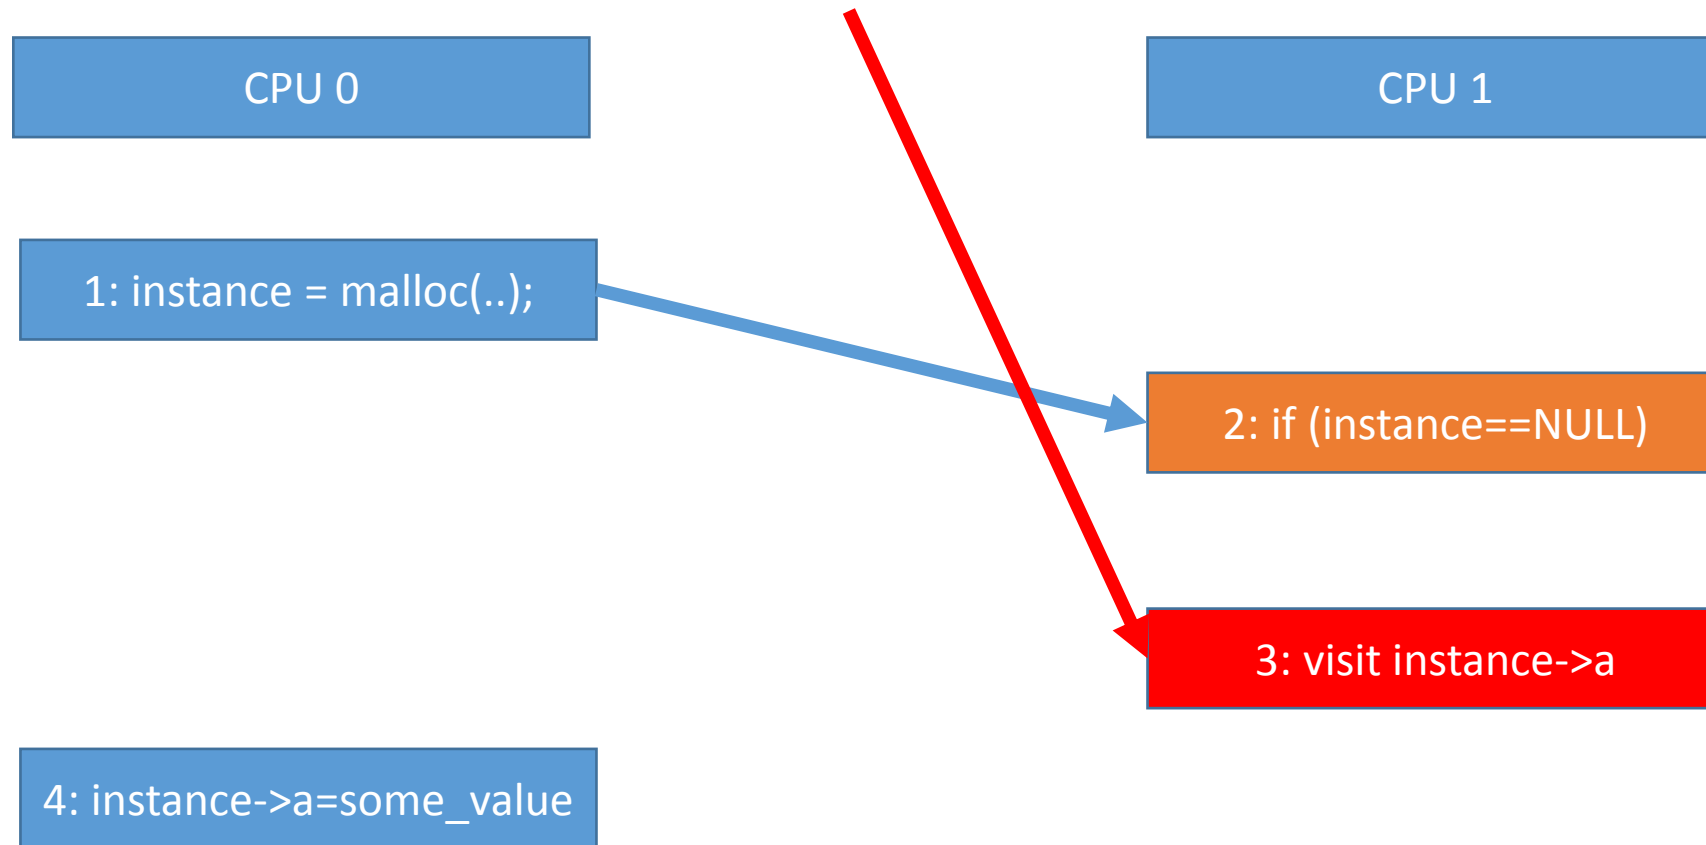
```
static struct T *instance;
static spinlock_t instance_lock;

struct T *get_instance()
{
    if (instance == NULL) {
        spin_lock(&instance_lock);
        if (instance == NULL) {
            instance = malloc(sizeof(struct T));
            instance->a = some_value;
        }
        spin_unlock(&instance_lock);
    }
    return instance;
}
```

Check first?

Won't work, Why?

A multi-thread-safe singleton



CPU 1 reads the uninitialized value of `->a`

A multi-thread-safe singleton: Use barrier()

```
static struct T *instance;
static spinlock_t instance_lock;

struct T *get_instance()
{
    struct T *tmp;
    if (instance == NULL) {
        spin_lock(&instance_lock);
        if (instance == NULL) {
            tmp = malloc(sizeof(struct T));
            tmp->a = some_value;

            barrier();

            instance = tmp;
        }
        spin_unlock(&instance_lock);
    }
}
```

Initialization before publication!

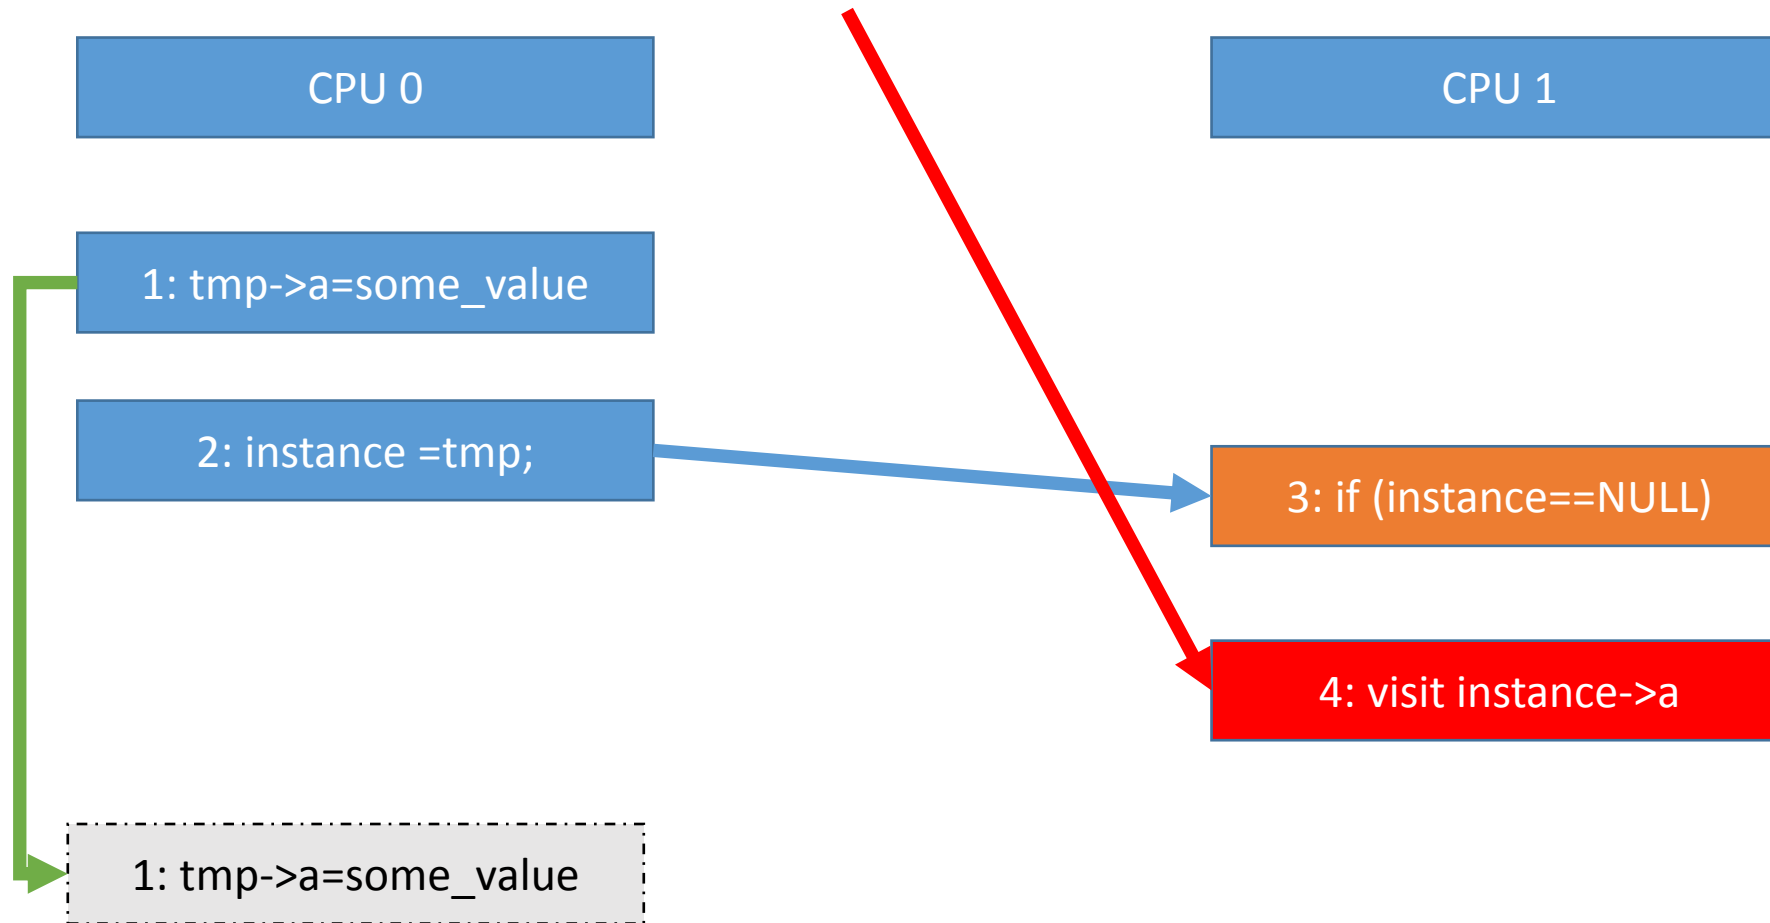
Why barrier() here?

barrier():

asm volatile(“” : : : “memory”);

Are we done?

A multi-thread-safe singleton: HW Reordering



CPU 1 still reads the uninitialized value of ->a

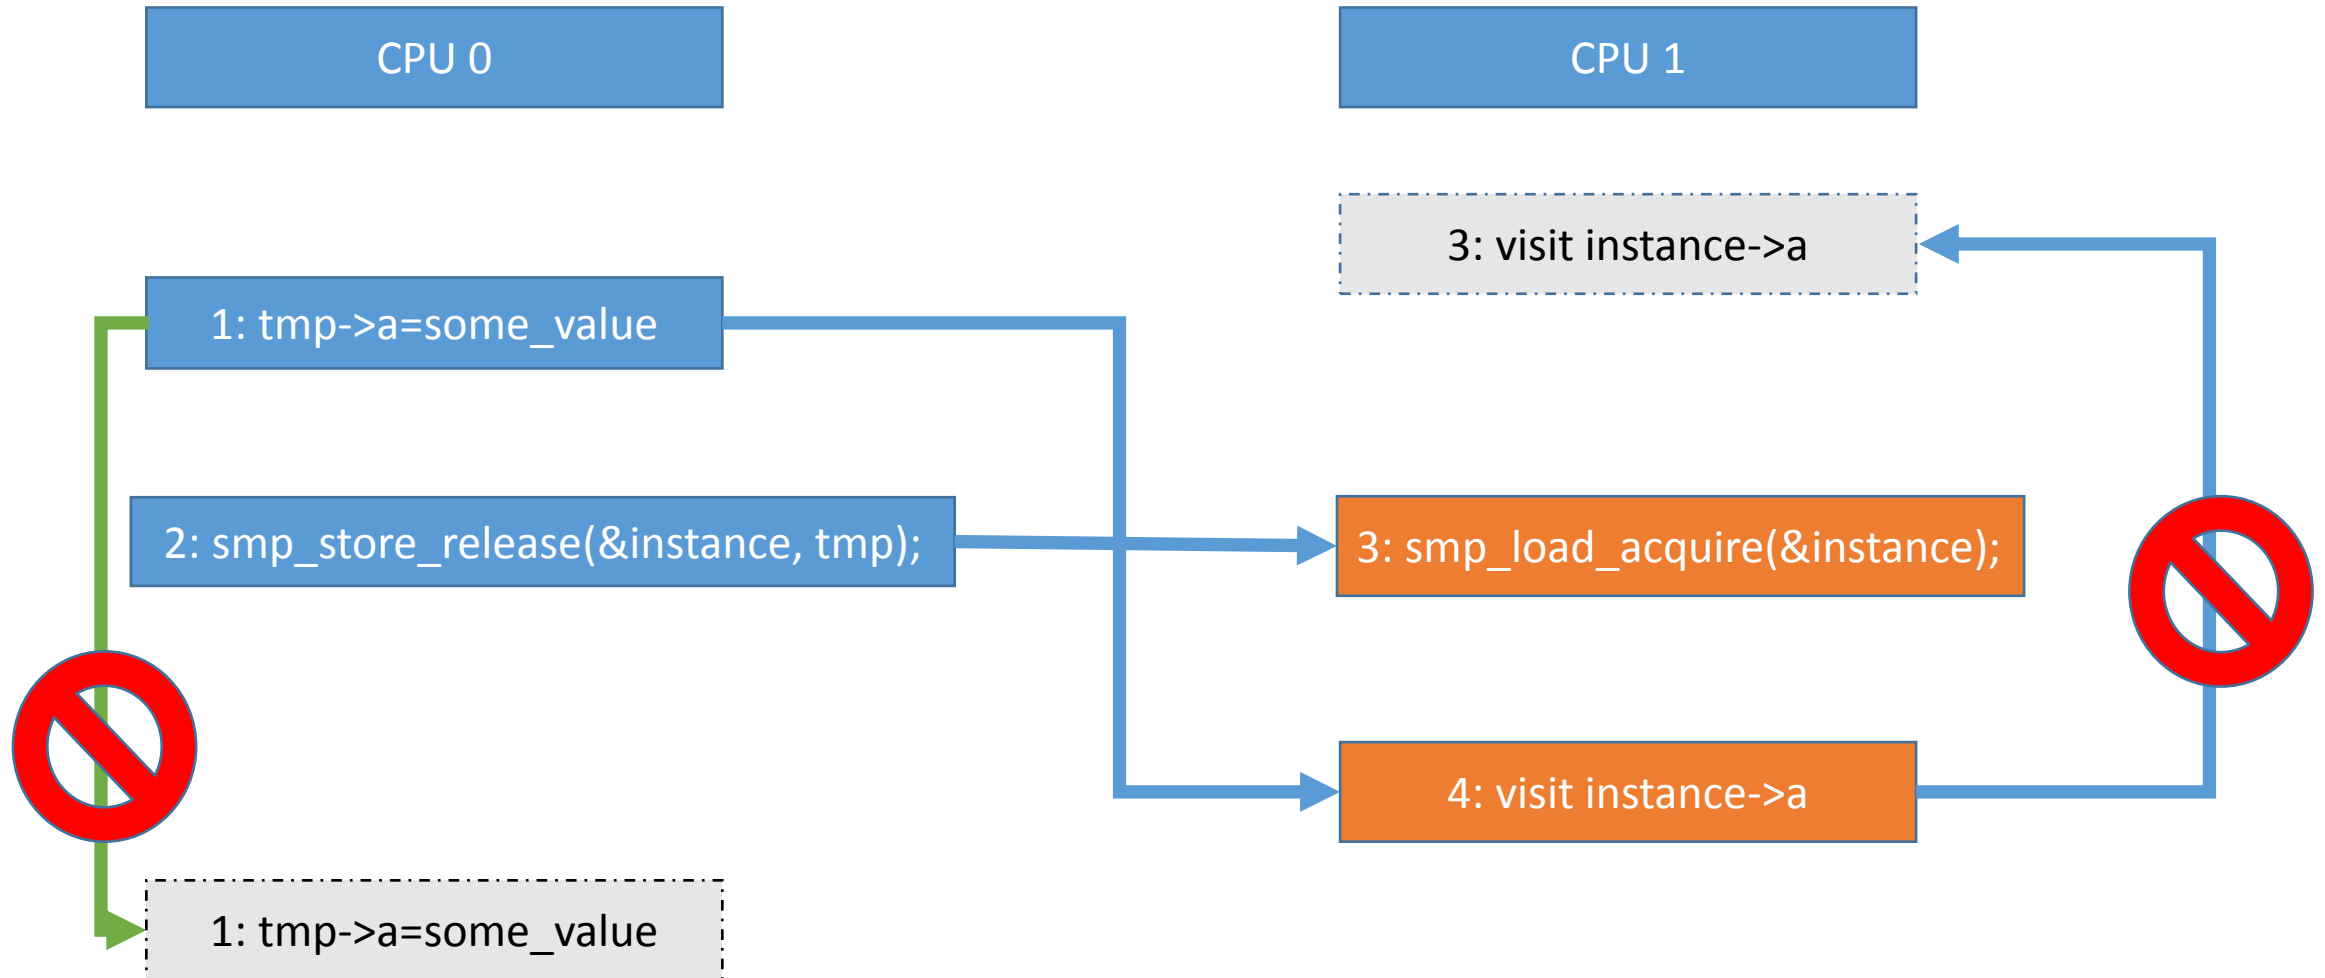
A multi-thread-safe singleton: Final Solution

Use acquire and release.

```
static struct T *instance;
static spinlock_t instance_lock;

struct T *get_instance()
{
    struct T *tmp = smp_load_acquire(&instance);
    if (tmp == NULL) {
        spin_lock(&instance_lock);
        if (instance == NULL) {
            tmp = malloc(sizeof(struct T));
            tmp->a = some_value;
            smp_store_release(&instance, tmp);
        }
        spin_unlock(&instance_lock);
    }
    return tmp;
}
```

Reordering is prohibited.



A multi-thread-safe singleton: Final Solution

Good enough?

```
static struct T *instance;
static spinlock_t instance_lock;

struct T *get_instance()
{
    struct T *tmp = smp_load_acquire(&instance);
    if (tmp == NULL) {
        spin_lock(&instance_lock);
        if (instance == NULL) {
            tmp = malloc(sizeof(struct T));
            tmp->a = some_value;
            smp_store_release(&instance, tmp);
        }
        spin_unlock(&instance_lock);
    }
    return tmp;
}
```