



# INTRODUCTION BLOCK MQ IO SCHEDULER

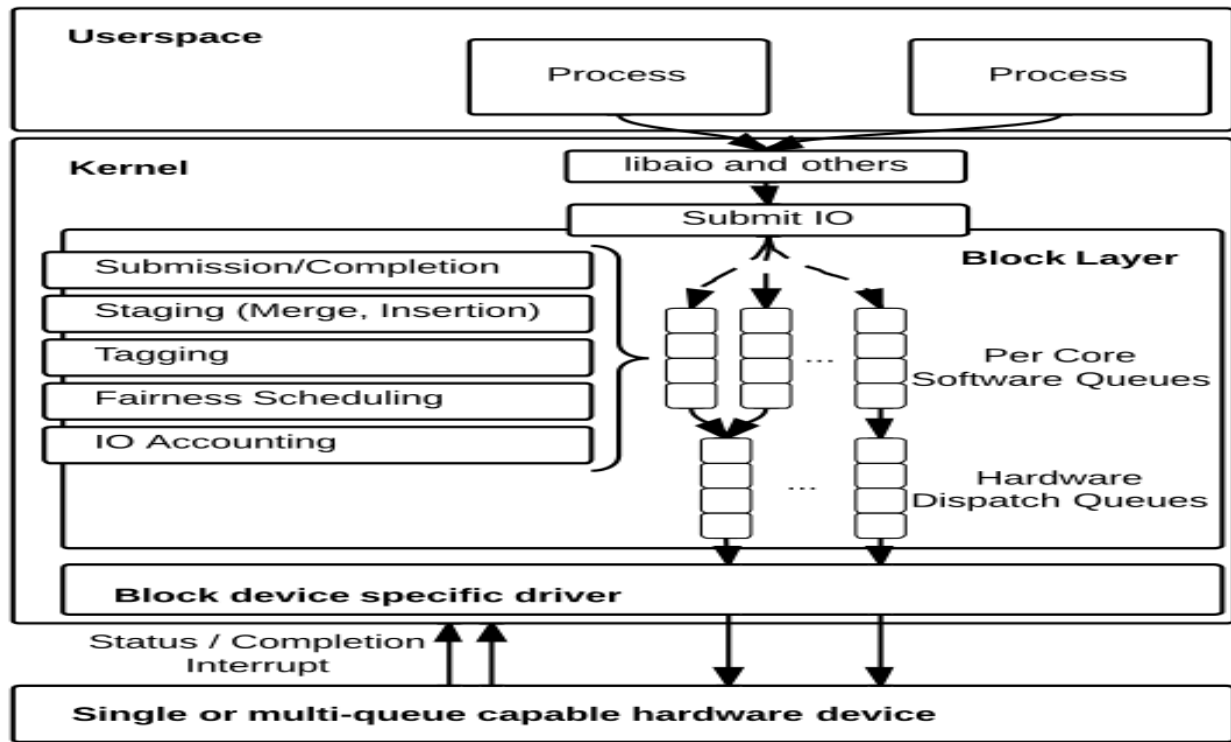
Ming Lei <[ming.lei@redhat.com](mailto:ming.lei@redhat.com)>, Red Hat

Oct. 22, 2017, Beijing, CLF2017

# Overview

- BLOCK MQ Background
  - Introduced in V3.13
  - for better supporting new storage of NVMe
  - address scalability issue of `q` → `queue_lock`
  - initially without any IO scheduler

# BLK MQ Framework

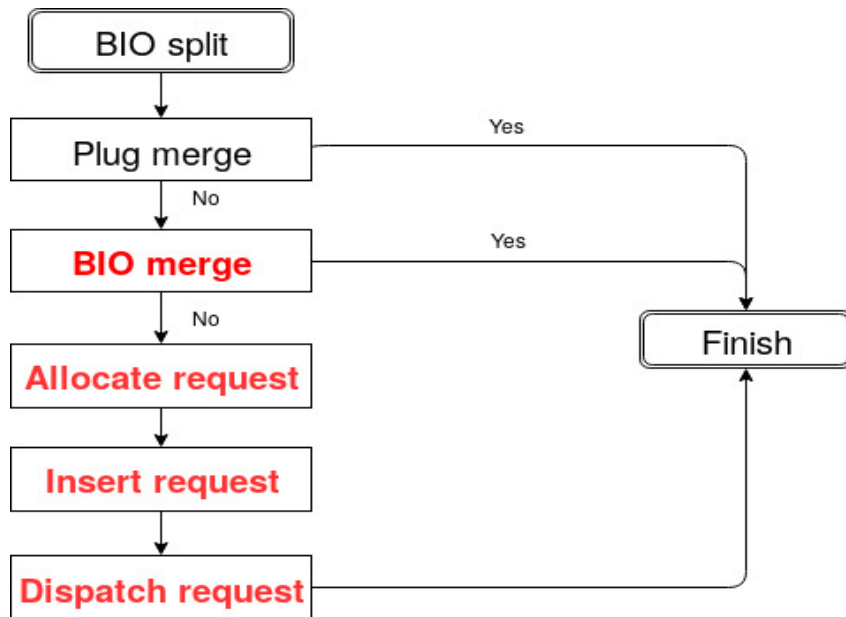


# BLOCK MQ IO SCHEDULER

- BLOCK MQ IO Scheduler background
- Better support traditional Storage device
- Replace blk\_queue\_bio() path totally in future
- Available in V4.11, initially with mq-deadline only
- BFQ and Kyber is merged in V4.12


# MQ IO Scheduler Framework

- BLOCK MQ IO Path(blk\_mq\_make\_request)



# MQ IO Merge Model

- IO scheduler queue Vs. LLD queue



IO Merge is possible IFF IO scheduler queue depth  $>$  LLD queue's

# MQ IO Merge Model

- IO scheduler queue depth
  - controlled via `/sys/block/XXX/queue/nr_requests`
  - respected via allocating request
- LLD queue depth
  - driver/device specific way to control, or not controllable
  - `.queue_rq()` returns `BLK_STS_RESOURCE` when LLD queue is full
  - LLD queue depth is highly related with storage device performance

# MQ IO Merge Model

- None Scheduler
  - introduced for NVMe at the beginning
  - no scheduler queue, so IO merge is possible IFF driver has specific queue depth, such as `q->queue_depth` on SCSI, not possible on NVMe actually
  - IO merge is on percpu SW queue, and use simple policy, merge isn't efficient



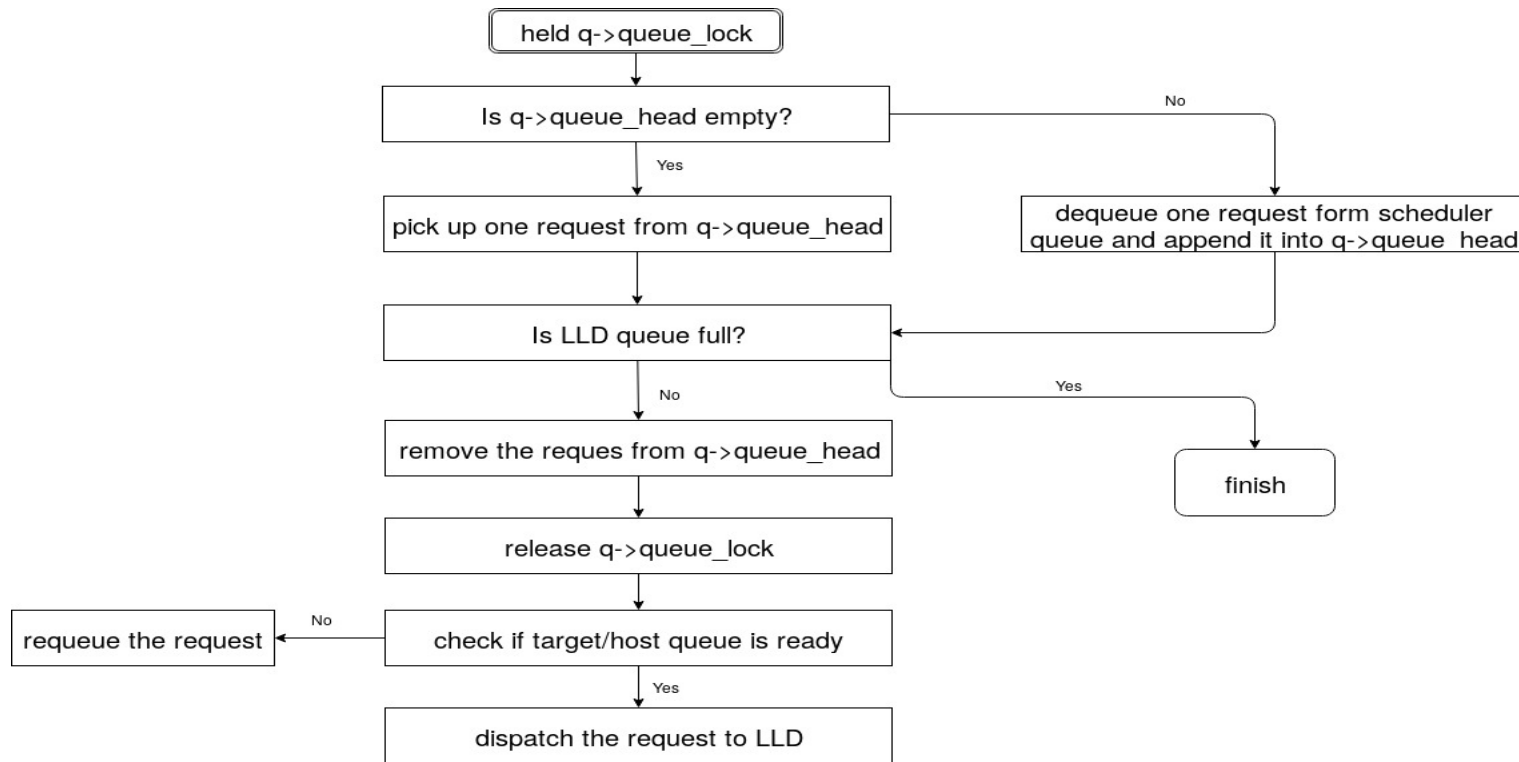
# MQ IO Merge Model

- MQ-DEADLINE / BFQ
  - basically similar with old block
  - introduced for making MQ working well on traditional disks (such as, SCSI)
  - IO merge is good because of per-request-queue scheduler queue
  - may not scale well for high performance MQ devices, such as NVMe, SCSI FC/SRP, because of per-request-queue lock

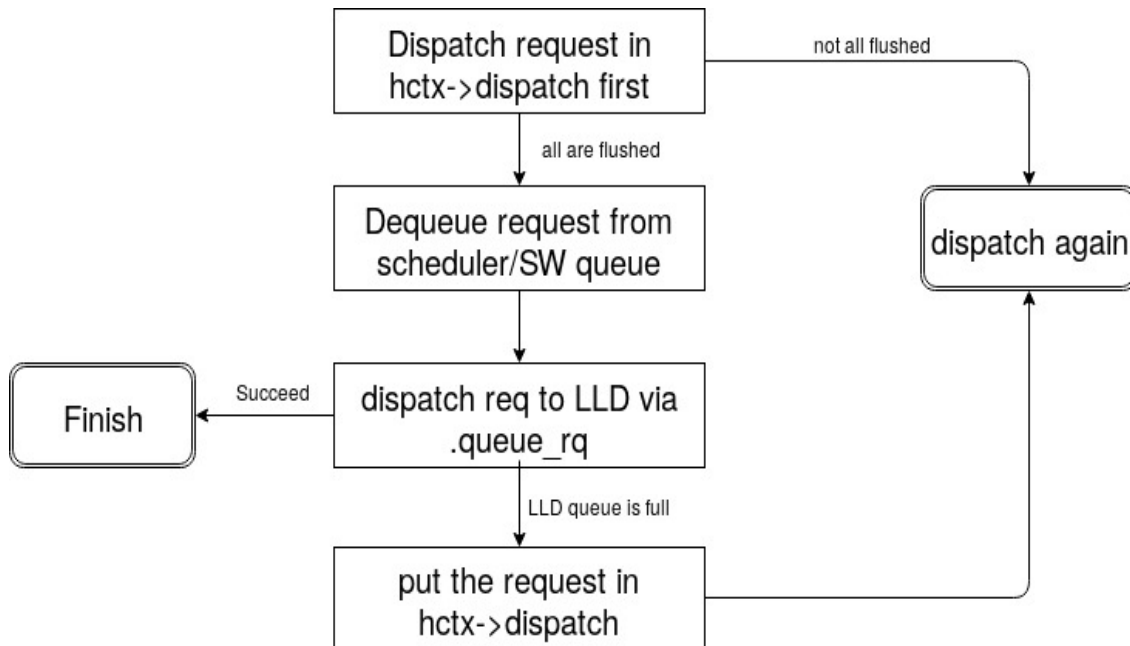
# MQ IO Merge Model

- Kyber
  - introduced for high performance devices, such as NVMe, NVMe OF
  - introduce READ, SYNC\_WRITE, OTHER domains, and each domain has its queue depth for simulating LLD queue depth,
  - IO merge is possible because of domain queue
  - IO merge is on percpu SW queue

# Block legacy IO Dispatch Model



# MQ IO Dispatch Model



# MQ IO Dispatch Model

- Issues
  - hctx->dispatch can't be dispatched one by one without holding hctx → lock; between moving hctx → dispatch moved to one temp list and being flushed out, scheduler can't be dequeued
  - q->queue\_depth is often among the whole request queue, all hctx should respect this limit

# MQ IO Dispatch Model

- Solutions for these issues
  - bypass hctx->dispatch totally
  - reserving budget before dequeuing from IO scheduler queue by introducing `.get_budget` and `.put_budget` in `blk_mq_ops`
  - will be merged to V4.15 if everything is fine
  - better than legacy path in theory without holding per-queue lock

# Performance data

- mq-deadline(fio, libaio, direct, bs=4k, queue\_depth=64, jobs=64, disk=SRP/IB, V4.14-rc4)

-----			
	V4.14-rc4	V4.14-rc4	patched V4.14-rc4
IOPS(K)	DEADLINE	MQ-DEADLINE	MQ-DEADLINE
-----			
read	450.0	154.12	474.0
-----			
write	419.65	135.88	481.89
-----			

# Next Step of MQ IO scheduler

- Improving on Kyber
  - pre-defined/hard coded domain depth
  - hard coded latency
  - domain queue depth adjust approach
  - very young
- SSD friendly IO schedule



# Next Step of MQ IO scheduler

- One big challenge
  - need to provide excellent support on modern high performance storage, such as NVMe, NVMe OF
  - meantime not cause performance regression on traditional storage, such as SCSI



# THANK YOU



[plus.google.com/+RedHat](https://plus.google.com/+RedHat)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[twitter.com/RedHatNews](https://twitter.com/RedHatNews)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)