



# INTEL<sup>®</sup> FPGA LINUX<sup>\*</sup> DRIVER SOLUTION AND UPSTREAMING STATUS

Figo Zhang, Hao Wu

# AGENDA

## Background

- Intel® FPGAs and the Modern Datacenter
- Platform Options and the Acceleration Stack
- Open Programmable Acceleration Engine

## Hardware Overview

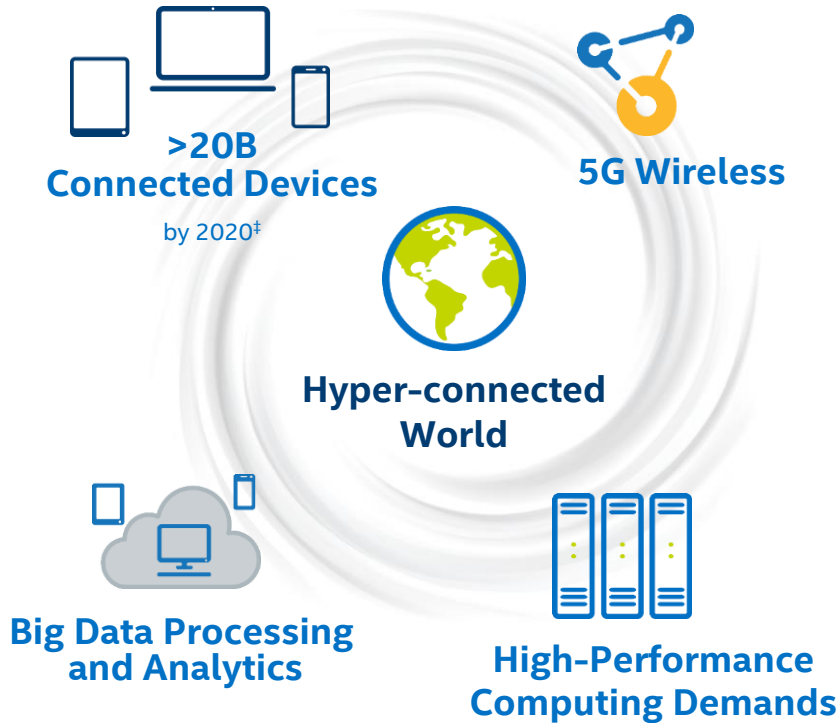
- FPGA Interface Manager (FIM)

## Intel FPGA Linux Driver solution

- Driver Architecture
- Partial Reconfiguration
- Virtualization
- Example: Simple DMA Operation
- Upstream Status


# Background

# DATA MOVEMENT AND PROCESSING EXPLOSION



## Markets

- Government
- Enterprise
- Cloud
- Communications



## Infrastructure

- Network
- Storage
- Compute



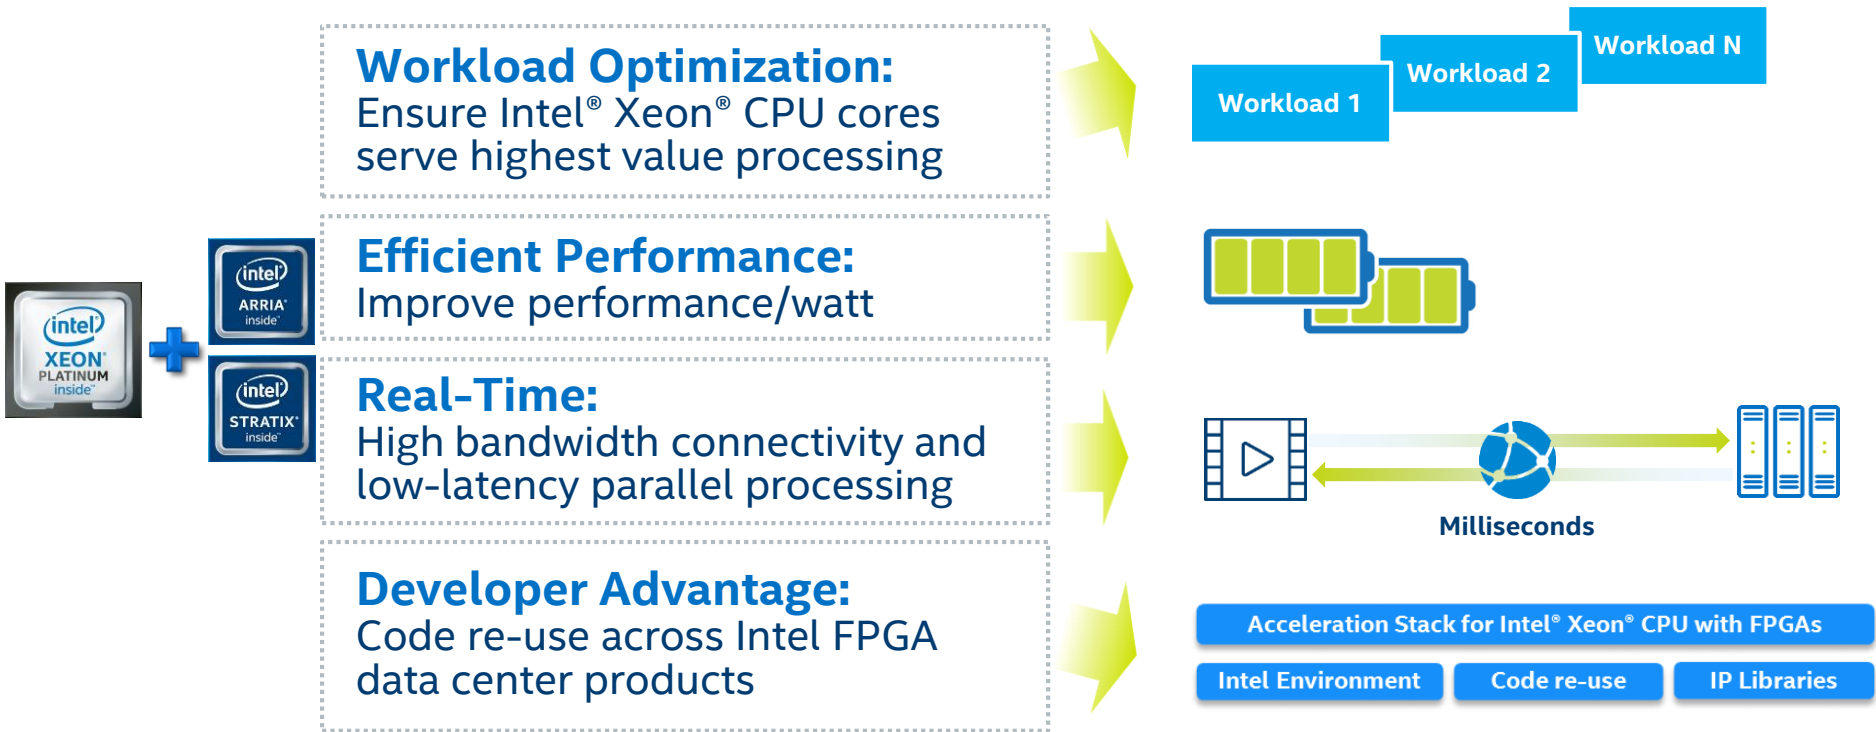
## Workloads

- Security
- Big Data Processing and Analytics
- Video processing and transcode
- Artificial Intelligence & Machine Learning
- Packet processing



<sup>†</sup> Source: "Gartner Says 8.4 Billion Connected 'Things' Will Be in Use in 2017, Up 31 Percent From 2016", 2/7/2017, <http://www.gartner.com/newsroom/id/3598917> (Table 1 - IoT Units Installed Base by Category, 2020 column - Grand Total, including consumer+business units)

# ACCELERATION ASSIST TO PROPEL DATA INSIGHT & OPERATIONAL EFFICIENCY




The Intel® Xeon® processor with FPGA acceleration can reduce TCO and solve new problems

# INTEL® FPGA DATA CENTER FORM FACTORS OPTIONS



Enabled By The Acceleration Stack for Intel® Xeon® CPU with FPGAs

## PCIe Acceleration Cards



- System flexibility with Intel Xeon CPU SKU options
- Dedicated local memory
- Can be slotted into 1U servers

## Server Platform Option with In-Package FPGA



- Coherent interface benefits software developers
- Superior performance for bandwidth & latency sensitive applications

Choose the Intel FPGA form factor matched to your application needs

# THE INTEL® APPLICATION DEVELOPER ADVANTAGE

Acceleration Stack for Intel® Xeon® CPU with FPGAs

Intel Environment

Code re-use

IP Libraries

## Acceleration Stack for Intel® Xeon® CPU with FPGAs – *Enhanced Performance, Simplified*

- Saves developer time to focus on unique value-add of their solution
- Enables unprecedented code re-use across multiple Intel FPGA form-factor products
- World's first common developer interface for Intel FPGA data center products
- Optimized and simplified hardware and software APIs provided by Intel
- Enables easier development and deployment of Intel FPGAs for workload optimization

The stable and optimized foundation for building your Intel FPGA-accelerated solution

# ACCELERATION STACK FOR INTEL® XEON® CPU WITH FPGAS

Enhanced Performance, Simplified

Dynamically Allocate Intel® FPGAs for Workload Optimization

**Rack-Level Solutions**



Simplified Application Development

**User Applications**

Deep Learning, Networking, Genomics, etc.

Leverage Common Frameworks

**Industry Standard SW Frameworks**



Fast-Track Your Performance

**Acceleration Libraries**

LZ4, Snappy, etc.

Workload Optimization with Less Effort

**Intel Developer Tools**

(Intel Parallel Studio XE, Intel FPGA SDK for OpenCL™, Intel Quartus® Prime)



Common Developer Interface for Intel FPGA Data Center Products

**Acceleration Environment**

(Intel Acceleration Engine with OPAE Technology, FPGA Interface Manager (FIM))




Intel® delivers a system-optimized solution stack for your data center workloads



# OPEN PROGRAMMABLE ACCELERATION ENGINE (OPAE) TECHNOLOGY

Simplified FPGA Programming Layer for Application Developers



**Consistent cross-platform API**

**Minimal software overhead and latency**

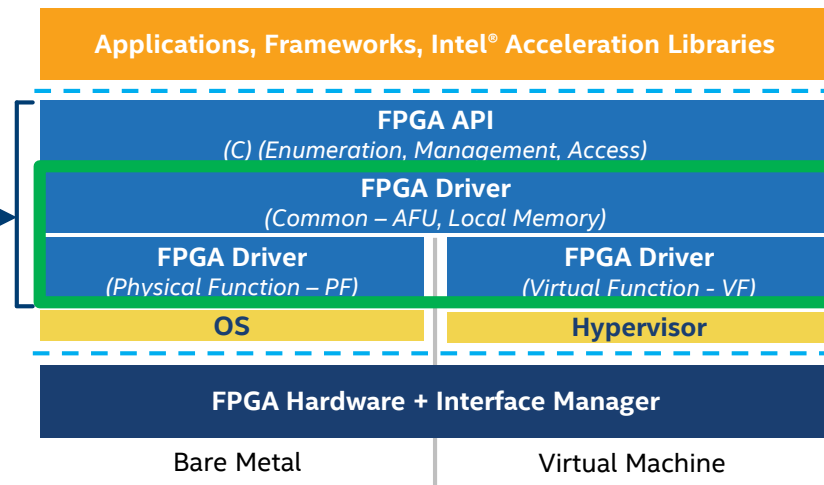
**Supports virtual machines and bare metal platforms**

**Open source code licensing and developer community** →

- Intel FPGA drivers being upstreaming to Linux kernel

SDK includes:

- Guides, utilities and sample code
- **AFU Simulation Environment (ASE)\*\*:**
  - Develop and debug Accelerator Functions faster

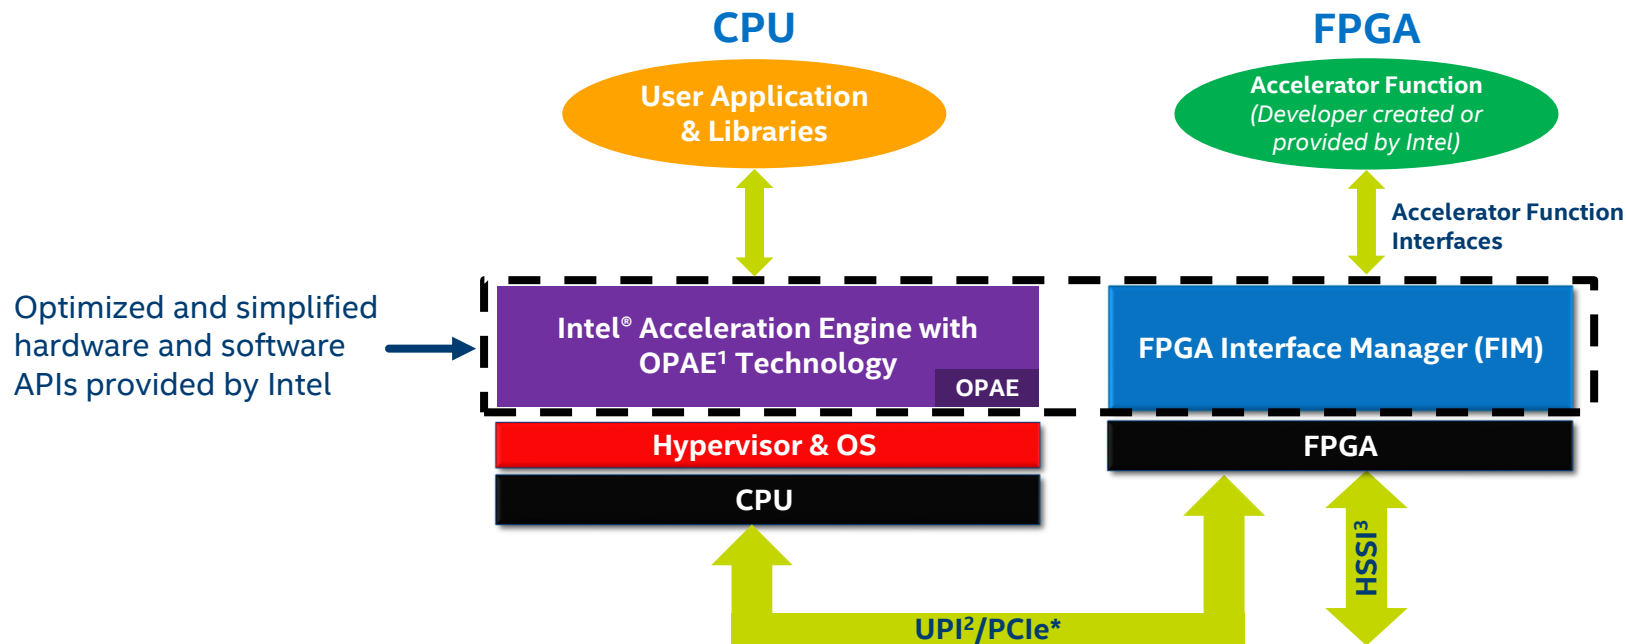



Start developing for Intel FPGAs with OPAAE today: <http://01.org/OPAAE>

# Hardware Overview

# ACCELERATION ENVIRONMENT

Common Developer Interface For Intel FPGA Data Center Products




<sup>1</sup>OPAE = Open Programmable Acceleration Engine

<sup>2</sup>UPI = Intel® Ultra Path Interconnect

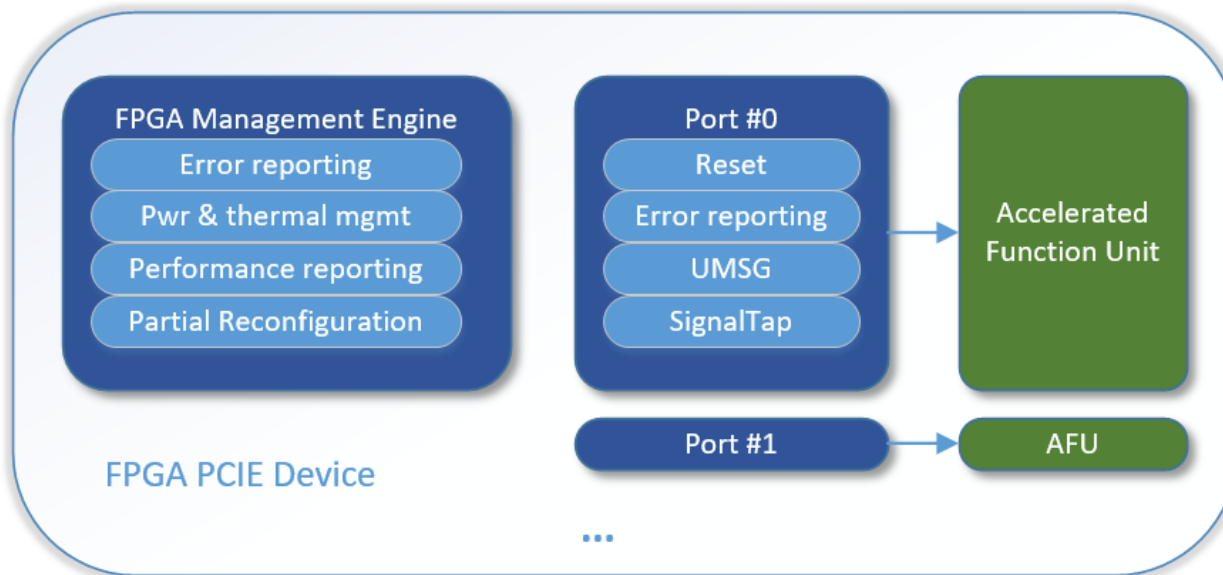
<sup>3</sup>HSSI = High Speed Serial Interface

# FPGA INTERFACE MANAGER (FIM) OVERVIEW




Device memory organized in Device Feature List data structure

Supported features exposed through Device Feature List



# FPGA INTERFACE MANAGER (FIM) DETAILS



- **FPGA Management Engine**

- Provides: power and thermal management, error reporting, partial reconfiguration, performance reporting, and other infrastructure functions.
- Each FPGA has one FME, accessible through the physical function.


- **Port**

- Interface between the static FPGA fabric (FIM) and a partially reconfigurable region containing an Accelerated Function Unit (Accelerator Function).
- Controls communication from SW and exposes features such as reset and debug.


- **Accelerated Function Unit**

- Attached to a port and exposes a MMIO region for accelerator-specific control registers.

# FPGA INTERFACE MANAGER (FIM) – VIRTUALIZATION SUPPORT

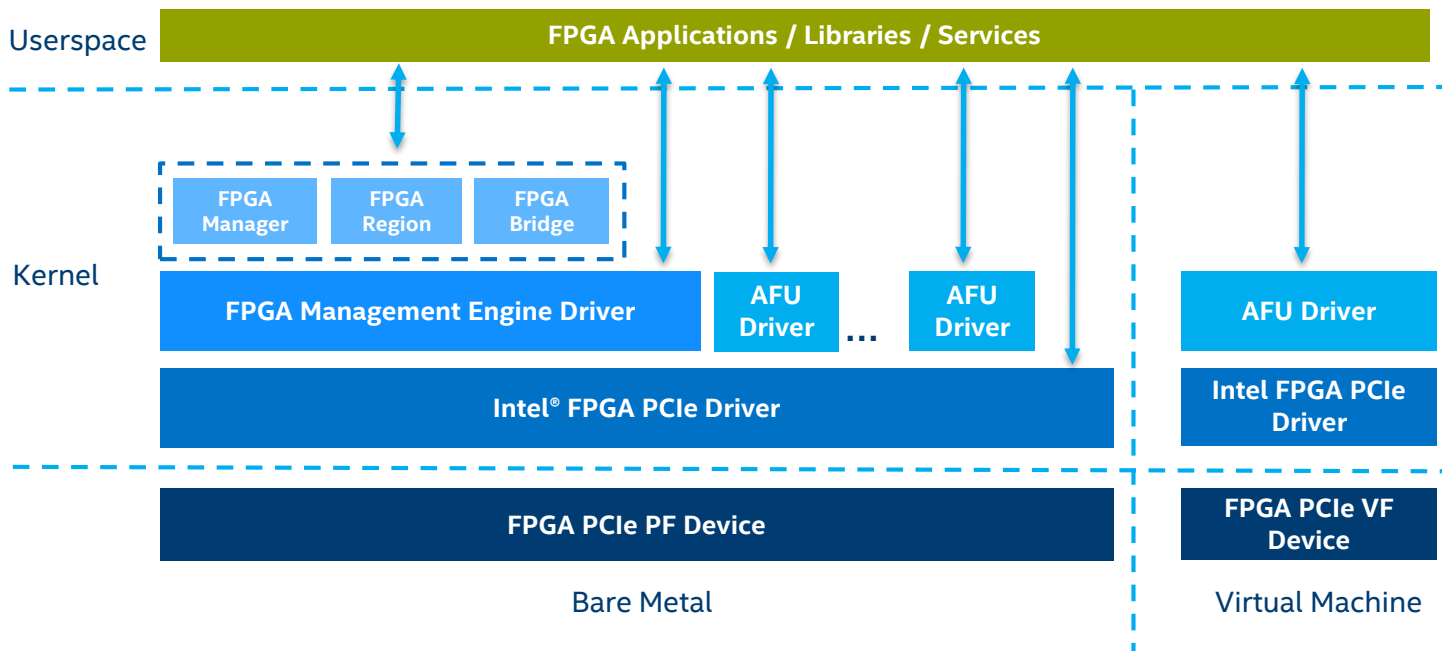


Supports PCIe SR-IOV function to create virtual functions (VFs) which can be used to assign individual accelerators to virtual machines.



# Intel® FPGA Linux Driver Solution

# INTEL® FPGA DRIVER ARCHITECTURE





# FPGA DRIVER COMPONENT – PCIE DEVICE DRIVER

- Enumeration
  - Discover Feature Devices by walking through the Device Feature List.
  - Create Platform Device with associated resources for Feature Devices.
    - Port/AFU and FME are Feature Devices.
  - Feature Device Framework:
    - Helper functions to manage feature devices and sub-features.
- SR-IOV Support
  - Sysfs interface to enable/disable VFs.
  - Both PF and VFs share the same driver.

# FPGA DRIVER COMPONENT – ACCELERATED FUNCTION UNIT

- Platform Device Driver
- Expose Accelerated Function Unit MMIO Resource
  - MMIO region (mmap)
- Provide DMA buffer mapping service
- Implement other sub features
  - Error reporting
  - UMSG (Unordered Message)
  - Debug

# ACCELERATED FUNCTION UNIT – DRIVER INTERFACES

## ▪ ioctl

- Get driver API version (FPGA\_GET\_API\_VERSION)
- Check for extensions (FPGA\_CHECK\_EXTENSION)
- Get port info (FPGA\_PORT\_GET\_INFO)
- Get MMIO region info (FPGA\_PORT\_GET\_REGION\_INFO)
- Map DMA buffer (FPGA\_PORT\_DMA\_MAP)
- Unmap DMA buffer (FPGA\_PORT\_DMA\_UNMAP)
- Reset AFU (FPGA\_PORT\_RESET)
- Enable UMsg (FPGA\_PORT\_UMSG\_ENABLE)
- Disable UMsg (FPGA\_PORT\_UMSG\_DISABLE)
- Set UMsg mode (FPGA\_PORT\_UMSG\_SET\_MODE)
- Set UMsg base address (FPGA\_PORT\_UMSG\_SET\_BASE\_ADDR)

## ▪ mmap

- mmap() accelerator MMIO regions.

## ▪ sysfs

- Path: /sys/class/fpga\_region/regionX/intel-fpga-afu.n/
- Read Accelerator GUID (afu\_id)
- Error Reporting (errors/)

# FPGA DRIVER COMPONENT – FPGA MANAGEMENT ENGINE

- Platform Device Driver
- Implements management features
  - FPGA capability and status.
  - Thermal & Power management.
  - Partial Reconfiguration function.
  - Global Error reporting.
  - Global Performance reporting.
  - Port Management.

# FPGA MANAGEMENT ENGINE – DRIVER INTERFACES

## ▪ ioctl

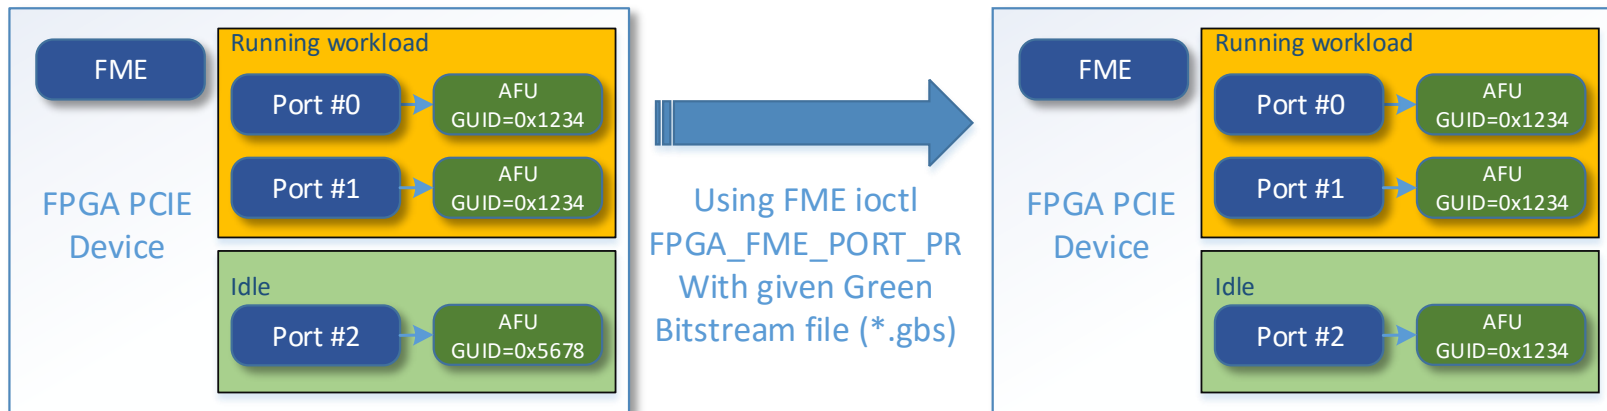
- Get driver API version (FPGA\_GET\_API\_VERSION)
- Check for extensions (FPGA\_CHECK\_EXTENSION)
- Assign port to PF (FPGA\_FME\_PORT\_ASSIGN)
- Release port from PF (FPGA\_FME\_PORT\_RELEASE)
- Program Bitstream (FPGA\_FME\_PORT\_PR)

## ▪ sysfs

- Path: /sys/class/fpga\_region/regionX/intel-fpga-fme.n/
- Read bitstream ID/metadata (bitstream\_id / bitstream\_metadata)
- Read number of ports (ports\_num)
- Read socket ID (socket\_id)
- Read performance counters (perf/)
- Power management (power\_mgmt/)
- Thermal management (thermal\_mgmt/)
- Error reporting (errors/)

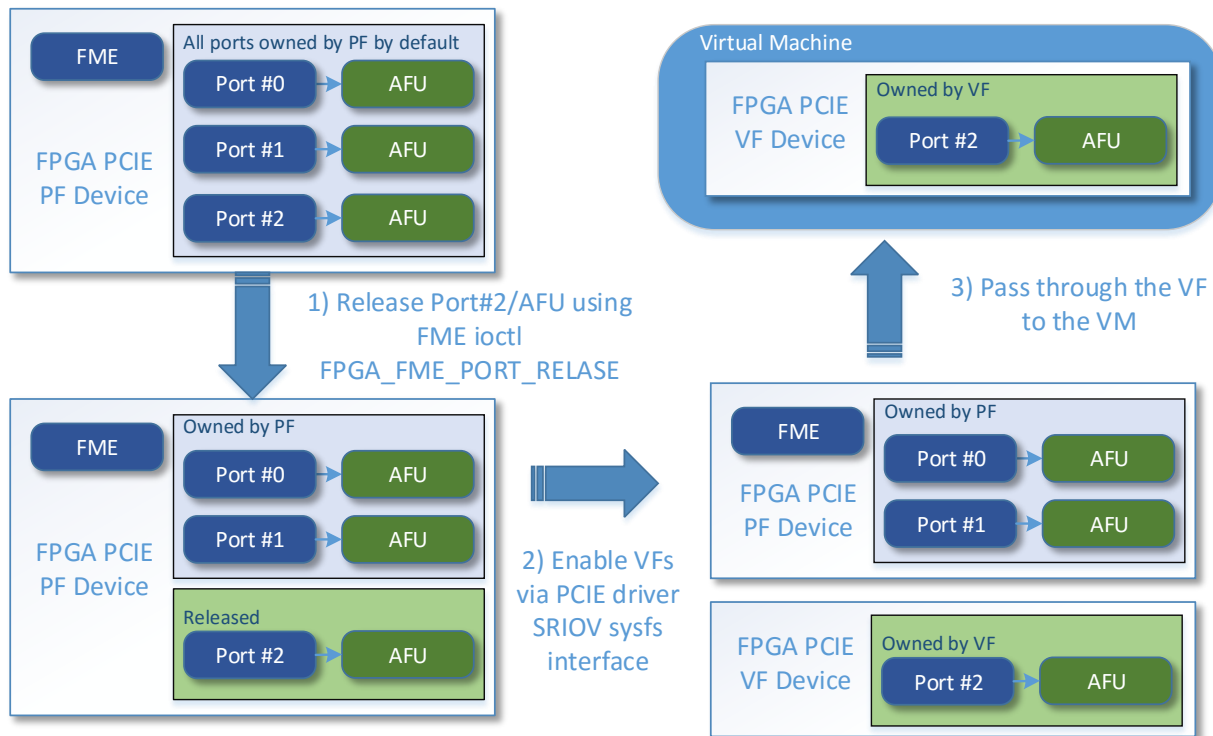
# PARTIAL RECONFIGURATION (PR)

- Accelerators reconfigured through partial reconfiguration.
- Other AFUs could run workload at the same time.
- Interface compatibility needed before start PR.



# VIRTUALIZATION

To enable access to an accelerator from within a virtual machine, AFU ports must be assigned to a VF



# EXAMPLE: SIMPLE DMA OPERATION

- Open AFU device file.
- Use AFU ioctl `FPGA_PORT_GET_REGION_INFO` to get AFU MMIO region information.
- Invoke AFU `mmap` to map AFU MMIO region for CSRs access.
- Allocate buffer and do DMA mapping via AFU ioctl `FPGA_PORT_DMA_MAP`.
- Program the DMA address to related CSRs in mapped area.
- Start DMA by programming related CSRs in mapped area.
- Poll on CSRs for completion.
- Unmap DMA by AFU ioctl `FPGA_PORT_DMA_UNMAP`
- Close AFU device file as task is done.



# UPSTREAMING STATUS

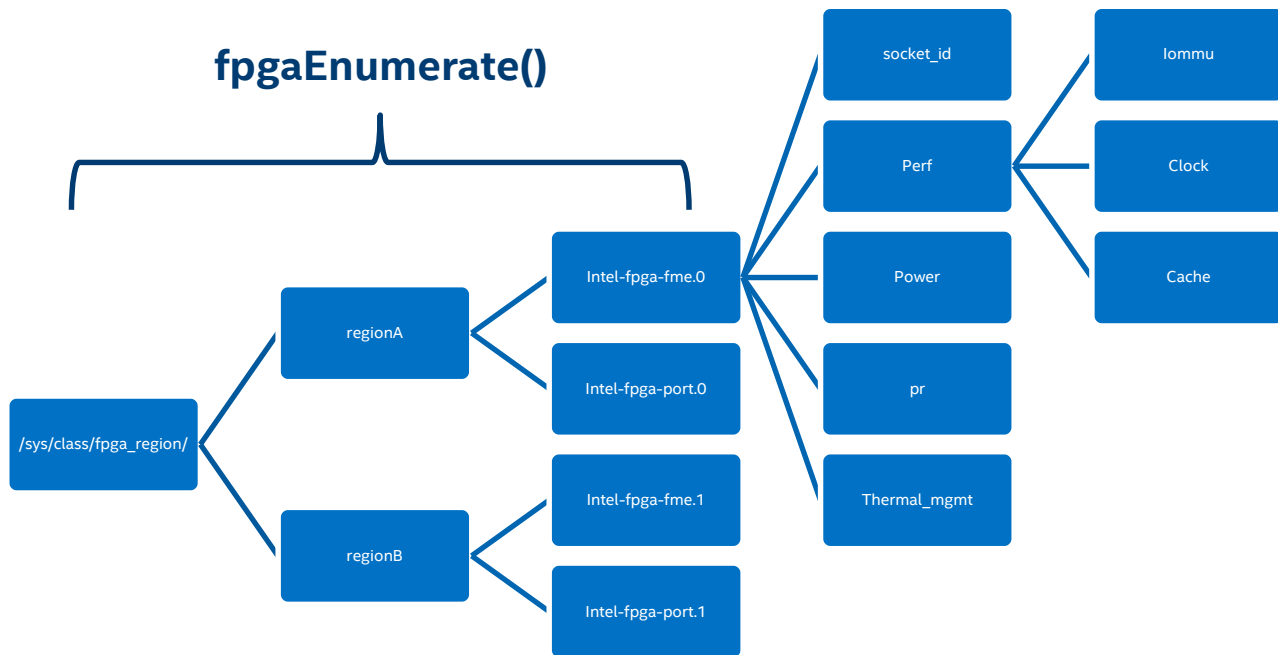
- First batch of patches has been submitted, version 2 is under review now.
  - Basic functions to enable AFU usage and Partial Reconfiguration.
    - Link: <https://marc.info/?l=linux-fpga&m=149844232609819&w=2>
  - Documentation/fpga/intel-fpga.txt:
    - <https://marc.info/?l=linux-fpga&m=149844234509825&w=2>
- More advanced features to be submitted next step
  - SR-IOV support and other sub features for FME and PORT/AFU.

# START DEVELOPING FOR INTEL® FPGAS WITH OPAE TODAY

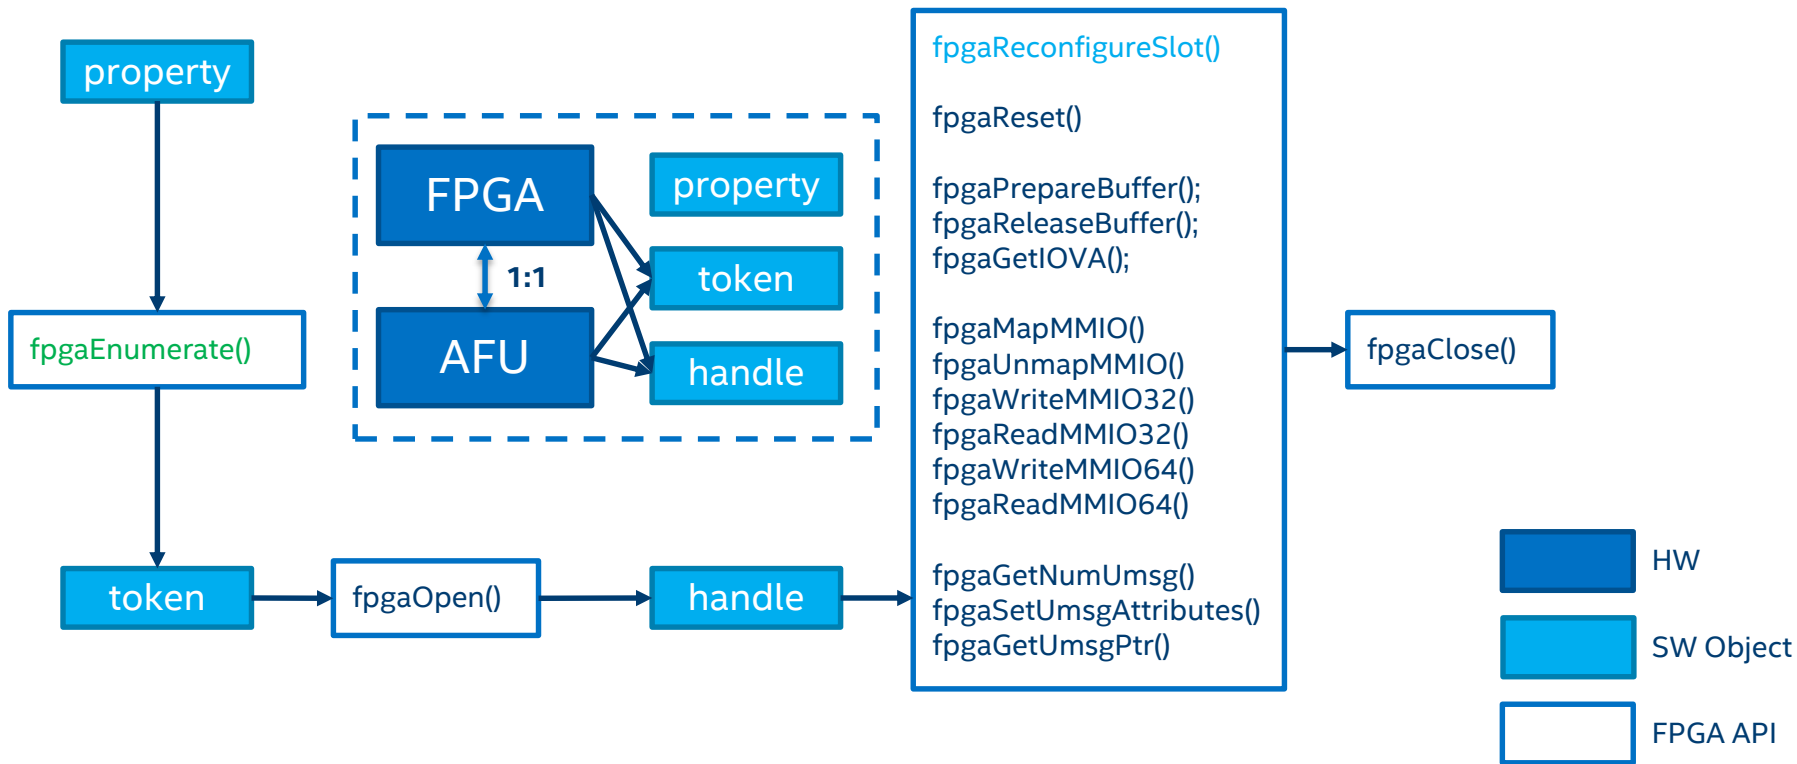
- Learn more about OPAE by visiting: <http://01.org/OPAE>
- Join the OPAE mailing list: <https://lists.01.org/pipermail/opae/>

**BACKUP**

# FPGA API - SYSFS & ENUMERATION



# FPGA API - ENUMERATE, MANAGE & ACCESS




# OPAE USERSPACE API ACCESS

```
fpga_result fpgaOpen(fpga_token token, fpga_handle *handle, int flags);
fpga_result fpgaClose(fpga_handle handle);
fpga_result fpgaReset(fpga_handle handle);
fpga_result fpgaPrepareBuffer(fpga_handle handle, uint64_t len, void **buf_addr, uint64_t *wsid, int
flags);
fpga_result fpgaReleaseBuffer(fpga_handle handle, uint64_t wsid);
fpga_result fpgaGetIOVA(fpga_handle handle, uint64_t wsid, uint64_t *iova);

fpga_result fpgaMapMMIO(fpga_handle handle, uint32_t mmio_num, uint64_t **mmio_ptr)
fpga_result fpgaUnmapMMIO(fpga_handle handle, uint32_t mmio_num)
fpga_result fpgaWriteMMIO32(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint32_t value)
fpga_result fpgaReadMMIO32(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint32_t *value)
fpga_result fpgaWriteMMIO64(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint64_t value)
fpga_result fpgaReadMMIO64(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint64_t *value)
fpga_result fpgaWriteMMIO(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint64_t value)
fpga_result fpgaReadMMIO(fpga_handle handle, uint32_t mmio_num, uint64_t offset, uint64_t *value)

fpga_result fpgaGetNumUmsg(fpga_handle handle, uint64_t *value)
fpga_result fpgaSetUmsgAttributes(fpga_handle handle, uint64_t value)
fpga_result fpgaGetUmsgPtr(fpga_handle handle, uint64_t **umsg_ptr)
```



# Notices & Disclaimers

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](http://intel.com).

No computer system can be absolutely secure.

Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit <http://www.intel.com/performance>.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit <http://www.intel.com/performance>.

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Cost reduction scenarios described are intended as examples of how a given Intel-based product, in the specified circumstances and configurations, may affect future costs and provide cost savings. Circumstances will vary. Intel does not guarantee any costs or cost reduction.

Intel does not control or audit third-party benchmark data or the web sites referenced in this document. You should visit the referenced web site and confirm whether referenced data are accurate.

© 2017 Intel Corporation. Intel, the Intel logo, Stratix, Arria, and Xeon are trademarks of Intel Corporation in the U.S. and/or other countries. \*Other names and brands may be claimed as property of others.