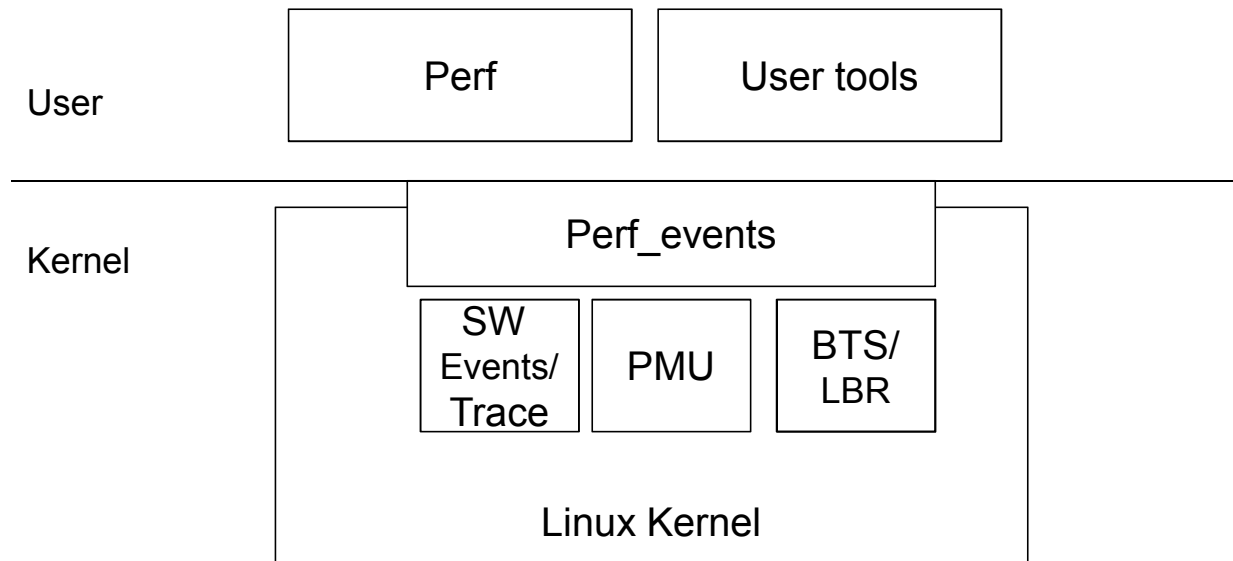


Application performance analysis

Using Perf with PMU event, PEBS, LBR and Intel PT technologies

Jin Yao (yao.jin@intel.com)

Linux “perf” overview



PMU, tracepoint, tracing framework

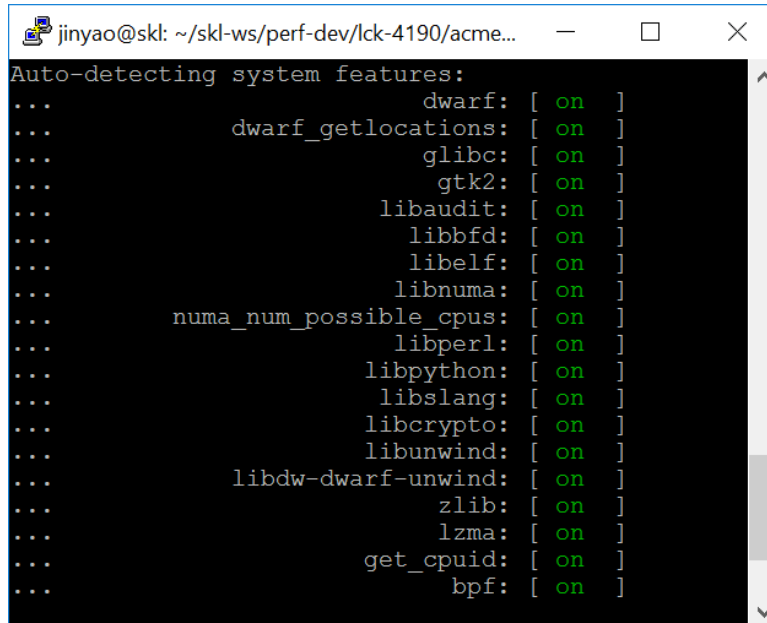
- Integrated into the Linux kernel
 - Including user tools
- Maintained by Linux community
 - With Intel contributions
- Generic: x86, other architectures
- Aims to abstract the hardware
- Supports software events
- Aims to be easy to use

Deployment

- Part of the core Linux kernel
- Fast development
- Not a separate driver
- Kernel version dependent, tightly integrated (some backports)
- Provides user interface (syscall + ring buffer)

Perf build notice

- Rebuild perf binary if use a new kernel
 - cd tools/perf; make
- Make sure the lib installed correctly



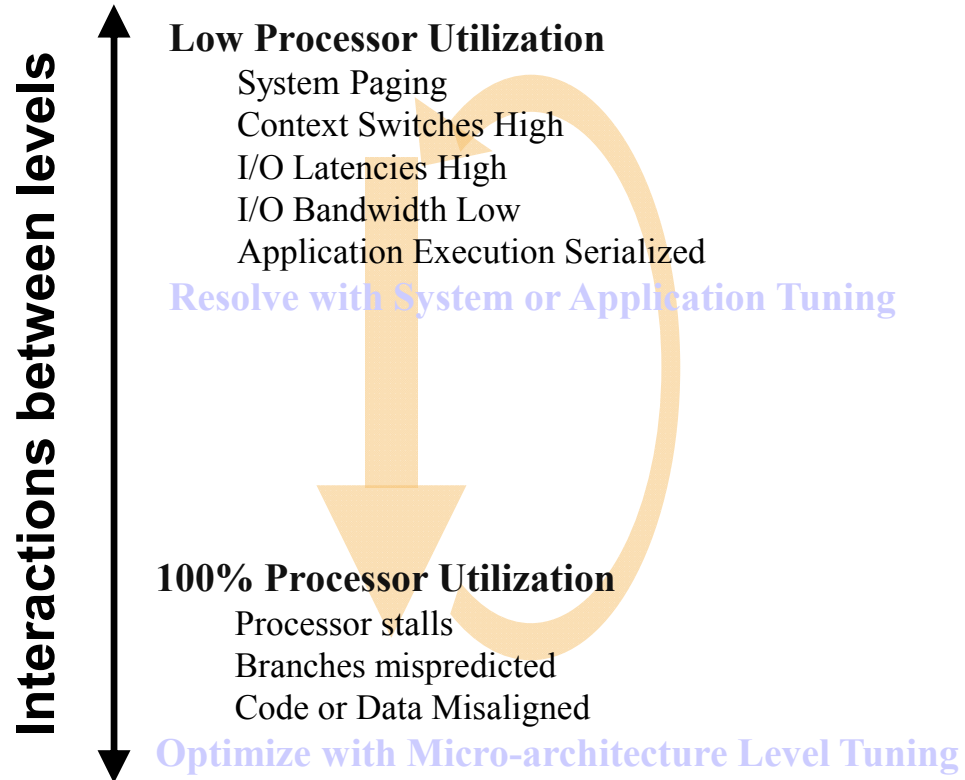
```
jinyao@skl: ~/skl-ws/perf-dev/lck-4190/acme...  
Auto-detecting system features:  
... dwarf: [ on ]  
... dwarf_getlocations: [ on ]  
... glibc: [ on ]  
... gtk2: [ on ]  
... libaudit: [ on ]  
... libbfd: [ on ]  
... libelf: [ on ]  
... libnuma: [ on ]  
... numa_num_possible_cpus: [ on ]  
... libperl: [ on ]  
... libpython: [ on ]  
... libslang: [ on ]  
... libcrypto: [ on ]  
... libunwind: [ on ]  
... libdw-dwarf-unwind: [ on ]  
... zlib: [ on ]  
... lzma: [ on ]  
... get_cpuid: [ on ]  
... bpf: [ on ]
```

A terminal window with a dark background and light text. The window title bar shows the user 'jinyao' at host 'skl' in the directory '~/skl-ws/perf-dev/lck-4190/acme...'. The terminal output shows 'Auto-detecting system features:' followed by a list of system features and their status, each on a new line. The status is '[on]' for all listed features. The features listed are: dwarf, dwarf_getlocations, glibc, gtk2, libaudit, libbfd, libelf, libnuma, numa_num_possible_cpus, libperl, libpython, libslang, libcrypto, libunwind, libdw-dwarf-unwind, zlib, lzma, get_cpuid, and bpf. The terminal has a scrollbar on the right side.

perf events I (perf list)

branch-instructions OR branches	[Hardware event]
branch-misses	[Hardware event]
bus-cycles	[Hardware event]
cache-misses	[Hardware event]
cache-references	[Hardware event]
cpu-cycles OR cycles	[Hardware event]
instructions	[Hardware event]
ref-cycles	[Hardware event]

Tuning Level Interactions



Data Collection Techniques

- Sampling
 - Collection of data based on the occurrence of a particular event such as a timer or interrupt
 - Example: Perf (perf record)
- Tracing
 - Getting log of path of application
 - Example: Perf (Intel PT)
- Instrumentation
 - Insertion of data collection instructions in the source code or object code level
- Simulation

A “mgen” workload example

- Generate Remote Memory Access for ~10s on SKX
 - `mgen -a 0 -c 28 -t 10` (memory allocated on node0, thread runs on cpu28)

Time	Latency (ns)
1.6s	157.5
3.2s	159.7
4.9s	158.6
6.5s	158.6
8.1s	158.6
9.7s	158.6
11.4s	158.6
Average	158.6

Overview (by perf stat)

- `perf stat -e cycles,instructions ./mgen -a 0 -c 28 -t 10`

```
Performance counter stats for './mgen -a 0 -c 28 -t 10':  
  
    35,069,650,829      cycles  
    399,206,040        instructions          #    0.01  insn per cycle  
  
    11.564120039 seconds time elapsed
```

- IPC = Instruction Per Cycle (**0.01, very bad data**)
- perf stat is not sampling

Who eats cycles? (by perf record/report)

- `perf record -e cycles ./mgen -a 0 -c 28 -t 10`
- `perf report --stdio` (`buf_read` eats 98.83% cycles)

```
#
# Overhead  Command  Shared Object  Symbol
# .....  ....
#
```

98.83%	mgen	mgen	[.] buf_read
0.53%	mgen	[kernel.kallsyms]	[k] clear_page_erms
0.38%	mgen	mgen	[.] rand_buf_init
0.08%	mgen	[kernel.kallsyms]	[k] clear_huge_page
0.02%	mgen	[kernel.kallsyms]	[k] _raw_spin_lock
0.02%	mgen	[kernel.kallsyms]	[k] _raw_spin_lock_irqsave
0.01%	mgen	[kernel.kallsyms]	[k] __irqentry_text_start
0.01%	mgen	[kernel.kallsyms]	[k] task_tick_fair
0.01%	mgen	[kernel.kallsyms]	[k] update_curr
0.01%	mgen	[kernel.kallsyms]	[k] __free_pages_ok
0.01%	mgen	mgen	[.] last_free_elem
0.00%	mgen	[kernel.kallsyms]	[k] account_user_time

- `perf record` is sampling.

Which instruction eats cycles? (by perf annotate)

- `perf annotate --stdio`

```
      :      void buf_read(void *buf, int read_num)
      :      {
0.00 :      417e74:      push    %rbp
0.00 :      417e75:      mov     %rsp,%rbp
0.00 :      417e78:      push    %rbx
0.00 :      417e79:      mov     %rdi,-0x10(%rbp)
0.00 :      417e7d:      mov     %esi,-0x14(%rbp)
      :      asm volatile (
0.00 :      417e80:      mov     $0x0,%esi
0.00 :      417e85:      mov     -0x10(%rbp),%rax
0.00 :      417e89:      mov     -0x14(%rbp),%ecx
0.00 :      417e8c:      mov     %esi,%ebx
0.00 :      417e8e:      mov     %rax,%rdx
0.00 :      417e91:      xor     %ebx,%ebx
      :
      :      0000000000417e93 <LOOP1>:
0.00 :      417e93:      mov     (%rdx),%rdx
99.97 :      417e96:      inc     %ebx
0.03 :      417e98:      cmp     %ecx,%ebx
0.00 :      417e9a:      jb      417e93 <LOOP1>
      :
```

- Is “inc %ebx” take 99.97% cycles in buf_read? **No!**

PEBS (Precise Event)

- no p - arbitrary skid
 - :p - constant skid
 - :pp - requested to have 0 skid (Intel PEBS events)
 - :ppp - must have 0 skid (only special case)
-
- Run perf record with precise option again
 - `perf record -e cycles:pp ./mgen -a 0 -c 28 -t 10`
 - If only perf record <app>, default is `-e cycles:ppp`

PEBS (Precise Event)

- `perf annotate --stdio`

```
      :      void buf_read(void *buf, int read_num)
      :      {
0.00 :      417e74:      push    %rbp
0.00 :      417e75:      mov     %rsp,%rbp
0.00 :      417e78:      push    %rbx
0.00 :      417e79:      mov     %rdi,-0x10(%rbp)
0.00 :      417e7d:      mov     %esi,-0x14(%rbp)
      :      asm volatile (
0.00 :      417e80:      mov     $0x0,%esi
0.00 :      417e85:      mov     -0x10(%rbp),%rax
0.00 :      417e89:      mov     -0x14(%rbp),%ecx
0.00 :      417e8c:      mov     %esi,%ebx
0.00 :      417e8e:      mov     %rax,%rdx
0.00 :      417e91:      xor     %ebx,%ebx
      :
      :      0000000000417e93 <LOOP1>:
99.69 :      417e93:      mov     (%rdx),%rdx
0.31 :      417e96:      inc     %ebx
0.00 :      417e98:      cmp     %ecx,%ebx
0.00 :      417e9a:      jnb     417e93 <LOOP1>
      :
      :      0000000000417e9c <STOP>:
      :      "cmp %2,%0\n\t"
```

- Why instruction at **417e93** takes 99.69% cycles in `buf_read`?

Memory load of 417e93 (by perf c2c)

- 99.69 : 417193: mov (%rdx), %rdx
- Why memory load so slow? Not hit in LLC? Not hit in local memory? Cache-line false-sharing issue?
- perf c2c record ./mgen -a 0 -c 28 -t 10
- perf c2c report --stdio
- c2c: cache to cache – Detect False-Sharing cache-lines.
- Based on Intel load latency facility.
 - Memory access of the access
 - Type of the access (e.g. remote memory hit?)
 - Latency (in cycles) of the load access

What's False-Sharing?

```
struct foo {
    int x;
    int y;
};

static struct foo f;

/* The two following functions are running concurrently: */

int sum_a(void)
{
    int s = 0;
    int i;
    for (i = 0; i < 10000000; ++i)
        s += f.x;
    return s;
}

void inc_b(void)
{
    int i;
    for (i = 0; i < 10000000; ++i)
        ++f.y;
}
```

sum_a re-read x from memory even though modification of y is irrelevant.

What data address hit by 417e93 (1)

- c2c can do more than False-Sharing analysis
- perf c2c report –stdio (part of output)

Trace Event Information		
Total records	:	39367
Locked Load/Store Operations	:	11
Load Operations	:	34296
Loads - uncacheable	:	1
Loads - IO	:	0
Loads - Miss	:	1
Loads - no mapping	:	0
Load Fill Buffer Hit	:	146
Load L1D hit	:	59
Load L2D hit	:	1
Load LLC hit	:	130
Load Local HITM	:	0
Load Remote HITM	:	248
Load Remote HIT	:	0
Load Local DRAM	:	5
Load Remote DRAM	:	33953
Load MESI State Exclusive	:	33953
Load MESI State Shared	:	5
Load LLC Misses	:	34206
LLC Misses to Local DRAM	:	0.0%
LLC Misses to Remote DRAM	:	99.3%
LLC Misses to Remote cache (HIT)	:	0.0%
LLC Misses to Remote cache (HITM)	:	0.7%

What data address hit by 417e93 (2)

- perf c2c report –stdio (part of output, actually many 417e93 entries)

#	Num	----- HITM ----- Rmt	----- Lcl ----- Lcl	-- Store Refs -- L1 Hit L1 Miss	Data address Offset	Pid	Code address	----- cycles ----- rmt hitm lcl hitm	load	Total records	cpu cnt	
#	0	1	0	0 0	0xfffff8e3f562e2a00							
#		100.00%	0.00%	0.00% 0.00%	0x30	13701	0xffffffff924da2ef	318 0	0	2	1	[k] acpi_ev_g
#	1	1	0	0 0	0x7f77c9f3cfc0							
#		100.00%	0.00%	0.00% 0.00%	0x0	13701	0x417e93	366 0	0	2	1	[.] buf_read
#	2	1	0	0 0	0x7f77ca6a8200							
#		100.00%	0.00%	0.00% 0.00%	0x0	13701	0x417e93	287 0	0	2	1	[.] buf_read
#	3	1	0	0 0	0x7f77cacc1200							
#		100.00%	0.00%	0.00% 0.00%	0x0	13701	0x417e93	323 0	0	2	1	[.] buf_read

- 417e93 generates a lot of remote memory access and almost no local or remote LLC hit (not false-sharing issue).

Timed LBR (Last Branch Records)

- Sampling + Tracing (h/w saves latest N branches to buffer)

```
* Given one basic block:
*
*   from      to      branch_i
*   * ----> *
*       |
*       | block
*       v
*       * ----> *
*       from    to      branch_i+1
*
* where the horizontal are the branches and the vertical is the executed
* block of instructions.
*
```

- Tell us the cycles of code block between 2 branches.

LBR sampling

Log LBRs at sample point
Support 32 entries on SKL

FROM	TO	CYCLES
123	456	5
...		

SKL/GLM ?						
	63	62	61	60:48	47:16	15:0
LBR_FROM_IP	SIGN_EXT (bit 47)				LBR FROM address	
LBR_TO_IP	SIGN_EXT (bit 47)				LBR TO address	
LBR_INFO	MISPRED	IN_TX	TSX_ABORTED	Reserved		cycle-count (*)

Sample using
Performance Counter



Cycles of hot code block

- `perf record -b -e cycles:pp ./mgen -a 0 -c 28 -t 10`
- `perf report --branch-history --stdio`

```
static
void buf_read(void *buf, int read_num)
{
    asm volatile (
        "xor %0, %0\n\t"
        "LOOP1:\n\t"
        "mov (%1),%1\n\t"
        "inc %0\n\t"
        "cmp %2,%0\n\t"
        "jb LOOP1\n\t"
        "STOP:\n\t"
        :: "b" (0), "d" (buf), "r" (read_num)
    );
}
```

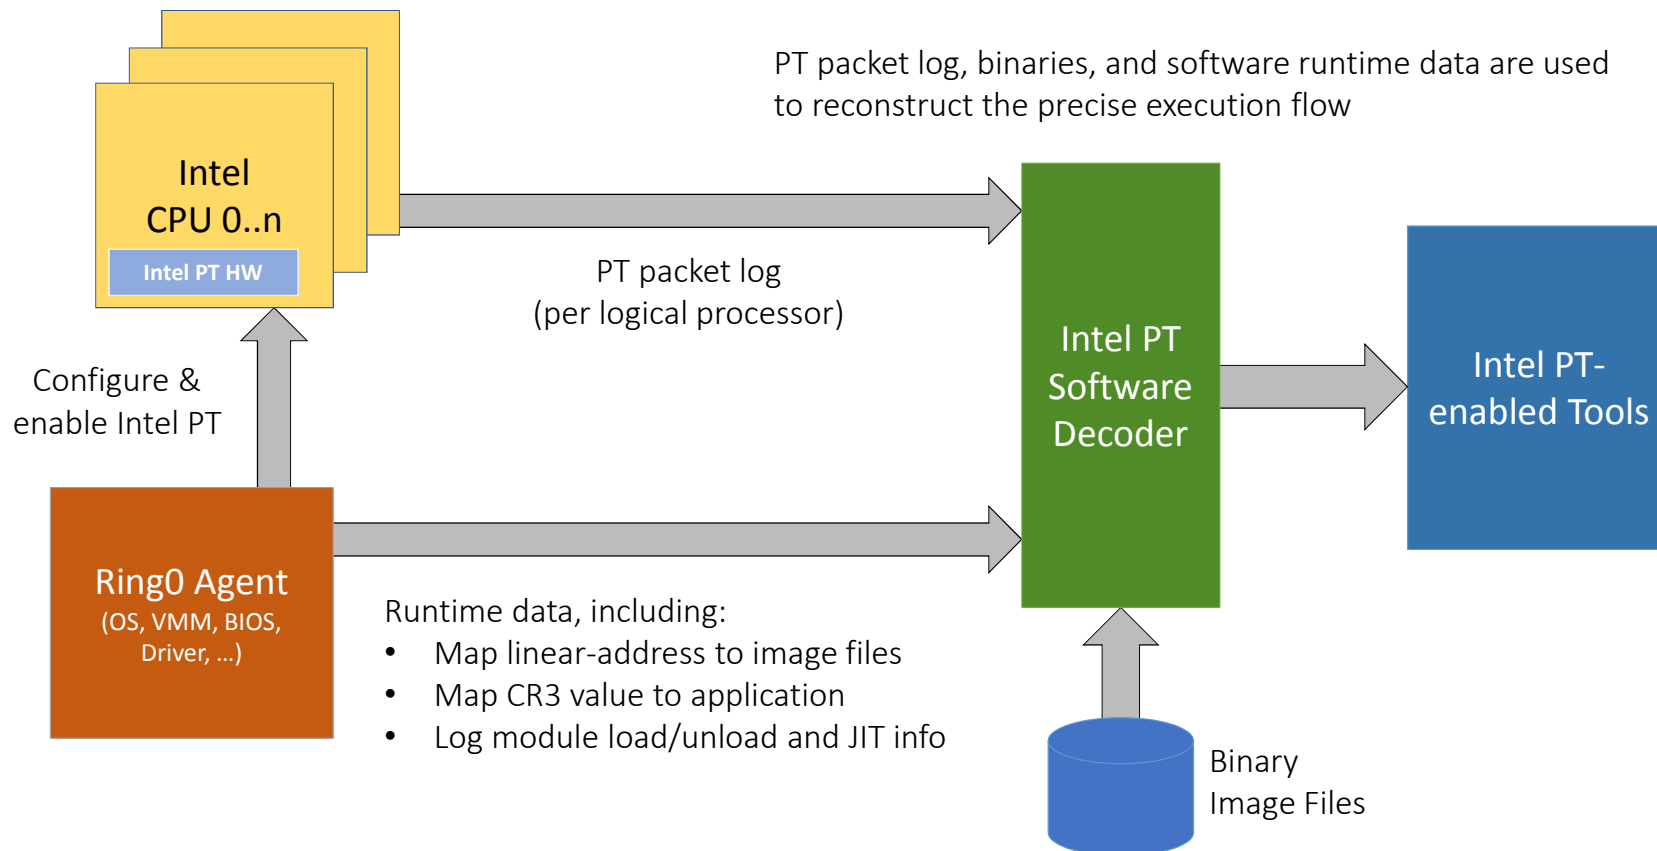
```
# Overhead Source:Line Symbol
# .....
#
# 98.44% util.c:38 [...] buf_read
# |
# --98.32%--buf_read util.c:38
# buf_read util.c:38 (cycles:440)
# buf_read util.c:38
# buf_read util.c:38 (cycles:440 iter:46044 avg_cycles:440)
# buf_read util.c:38
# buf_read util.c:38 (cycles:439 iter:92088 avg_cycles:439)
# buf_read util.c:38
# buf_read util.c:38 (cycles:439 iter:138132 avg_cycles:440)
# buf_read util.c:38
# buf_read util.c:38 (cycles:439 iter:184176 avg_cycles:440)
# buf_read util.c:38
# buf_read util.c:38 (cycles:443 iter:230220 avg_cycles:441)
# buf_read util.c:38
# buf_read util.c:38 (cycles:440 iter:276264 avg_cycles:441)
```

- Yellow is TO of branch X (LOOP1), green is FROM of branch X+1 (jb)
- 440 cycles is for code block from LOOP1 to jb

What is Processor Trace (PT)?

- Intel PT is a hardware feature that logs information about software execution
- Available in Skylake, Goldmont, ... Broadwell also, but has many limitations and is slower
- Supports control flow tracing. Decoder can determine exact flow of software execution from trace log
 - Target <5% performance overhead. Depends on processor generation and usage model
- Can store both cycle count and timestamp information

Intel® Processor Trace Components



Branch timestamp (by perf PT)

- `perf record -e intel_pt//u ./mgen -a 0 -c 28 -t 10`
- `perf script --ns -F time,cpu,sym,ip,srcline`

```
0000000000417e93 <LOOP1>:
417e93:    48 8b 12          mov     (%rdx),%rdx
417e96:    ff c3           inc     %ebx
417e98:    39 cb           cmp     %ecx,%ebx
417e9a:    72 f7           jb      417e93 <LOOP1>
```

```
[028] 427634.414462047:      417e9a buf_read
      util.c:38
[028] 427634.414462367:      417e9a buf_read
      util.c:38
[028] 427634.414462687:      417e9a buf_read
      util.c:38
[028] 427634.414463007:      417e9a buf_read
      util.c:38
[028] 427634.414463327:      417e9a buf_read
      util.c:38
[028] 427634.414463647:      417e9a buf_read
      util.c:38
[028] 427634.414463967:      417e9a buf_read
      util.c:38
```


Other Tools - NumaTOP

- NumaTOP (runtime memory locality characterization on NUMA system)

```
NumaTOP v2.0, (C) 2015 Intel Corporation
Monitoring 1047 processes and 1196 threads (interval: 5.0s)

  PID      PROC      RMA (K)    LMA (K)    RMA/LMA    CPI    *CPU%
17586      mgen      31654.4     7.1       4467.8     62.20    0.9
17577      numatop      15.6       38.0        0.4       1.11     0.0
4948      irqbalance      1.4        1.1        1.3       0.57     0.0
  1        systemd      0.0        0.0        0.0       0.00     0.0
  2        kthreadd      0.0        0.0        0.0       0.00     0.0
  3        kworker/0:0      0.0        0.0        0.0       0.00     0.0
  4        kworker/0:0      0.0        0.0        0.0       0.00     0.0
  6        kworker/u67      0.0        0.0        0.0       0.00     0.0
  7        mm_percpu_w      0.0        0.0        0.0       0.00     0.0
  8        ksoftirqd/0      0.0        0.0        0.0       0.00     0.0
  9        rcu_sched      0.0        0.0        0.0       0.00     0.0
 10        rcu_bh      0.0        0.0        0.0       0.00     0.0
 11        migration/0      0.0        0.0        0.0       0.00     0.0
 12        watchdog/0      0.0        0.0        0.0       0.00     0.0

<- Hotkey for sorting: 1(RMA), 2(LMA), 3(RMA/LMA), 4(CPI), 5(CPU%) ->
CPU% = system CPU utilization

Q: Quit; H: Home; R: Refresh; I: IR Normalize; N: Node
```

- <http://01.org/numatop>
- <https://github.com/01org/numatop.git>

Other Tools – LKP-tests

- LKP-tests (Linux kernel performance test tool)
- Open source tool by Intel:

<https://github.com/01org/lkp-tests.git>

- Framework to run benchmarks

Integrated ~80 benchmarks/test suites

Flexible mechanism to configure various parameters

Integrated ~40 monitors to monitor resource usages and statistics

- Framework for performance analysis
- Can be set up in CI environment (e.g. 0-Day CI), used for running benchmark and reproducing regression

References

Perf C2C:

<https://joemario.github.io/blog/2016/09/01/c2c-blog/>

LBR doc:

<http://lwn.net/Articles/680985/>

<http://lwn.net/Articles/680996/>

Perf PT doc:

<https://git.kernel.org/cgit/linux/kernel/git/torvalds/linux.git/tree/tools/perf/Documentation/intel-pt.txt>

Adding processor trace to Linux

<https://lwn.net/Articles/648154/>